

From Security Protocols to Pushdown Automata

RÉMY CHRÉTIEN, LORIA & LSV - CNRS, France
 VÉRONIQUE CORTIER, LORIA - CNRS, France
 STÉPHANIE DELAUNE, LSV, ENS Cachan & CNRS, France

Formal methods have been very successful in analyzing security protocols for reachability properties such as secrecy or authentication. In contrast, there are very few results for equivalence-based properties, crucial for studying, for example, privacy-like properties such as anonymity or vote secrecy.

We study the problem of checking equivalence of security protocols for an unbounded number of sessions. Since replication leads very quickly to undecidability (even in the simple case of secrecy), we focus on a limited fragment of protocols (standard primitives but pairs, one variable per protocol's rules) for which the secrecy preservation problem is known to be decidable. Surprisingly, this fragment turns out to be undecidable for equivalence. Then, restricting our attention to deterministic protocols, we propose the first decidability result for checking equivalence of protocols for an unbounded number of sessions. This result is obtained through a characterization of equivalence of protocols in terms of equality of languages of (generalized, real-time) deterministic pushdown automata. We further show that checking for equivalence of protocols is actually equivalent to checking for equivalence of generalized, real-time deterministic pushdown automata.

Very recently, the algorithm for checking for equivalence of deterministic pushdown automata has been implemented. We have implemented our translation from protocols to pushdown automata, yielding the first tool that decides equivalence of (some class of) protocols, for an unbounded number of sessions. As an application, we have analyzed some protocols of the literature including a simplified version of the basic access control (BAC) protocol used in biometric passports.

CCS Concepts: • **Security and privacy** → **Formal methods and theory of security**; *Logic and verification*

Additional Key Words and Phrases: Formal proofs, security protocols, verification, trace equivalence

ACM Reference Format:

Rémy Chrétien, Véronique Cortier, and Stéphanie Delaune. 2015. From security protocols to pushdown automata. *ACM Trans. Comput. Logic* 17, 1, Article 3 (September 2015), 45 pages.
 DOI: <http://dx.doi.org/10.1145/2811262>

1. INTRODUCTION

Formal methods have been successfully applied for rigorously analyzing security protocols. In particular, many algorithms and tools (Rusinowitch and Turuani [2003], Blanchet [2001], Comon-Lundh and Cortier [2003], Basin et al. [2005], and Cremers [2008], to cite a few) have been designed to automatically find flaws in protocols or prove

The research leading to these results has received funding from the European Research Council under the European Union's Seventh Framework Programme (FP7/2007-2013) / ERC grant agreement n° 258865, project ProSecure, and the ANR project JCJC VIP n° 11 JS02 006 01.

Authors' addresses: R. Chrétien and S. Delaune, Laboratoire Spécification et Vérification, 61, avenue du président Wilson, 94230 Cachan - France; emails: {chretien, delaune}@sv.ens-cachan.fr; V. Cortier, Laboratoire lorrain de recherche en informatique et ses applications, Campus Scientifique, BP 239, 54506 Vandoeuvre-lès-Nancy Cedex France; email: cortier@loria.fr.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or permissions@acm.org.

© 2015 ACM 1529-3785/2015/09-ART3 \$15.00

DOI: <http://dx.doi.org/10.1145/2811262>

security. Most of these results focus on reachability properties such as authentication or secrecy: for any execution of the protocol, it should never be the case that an attacker learns some secret (secrecy property) or that an attacker makes Alice think she's talking to Bob while Bob did not engage in a conversation with her (authentication property). However, privacy properties such as vote secrecy, anonymity, or untraceability cannot be expressed as reachability properties. They are instead defined as indistinguishability properties in Arapinis et al. [2010] and Bruso et al. [2010]. For example, Alice's identity remains private if an attacker cannot distinguish a session where Alice is talking from a session where Bob is talking.

Studying indistinguishability properties for security protocols amounts to checking a behavioral equivalence between processes. Processes represent protocols and are specified in some process algebras such as CSP or pi-calculus, except that messages are no longer atomic actions but terms, in order to faithfully represent cryptographic messages. Of course, considering terms instead of atomic actions considerably increases the difficulty of checking equivalence. As a matter of fact, there are just a few results for checking equivalence of processes that manipulate terms.

- Based on a procedure developed in Baudet [2005], it has been shown that trace equivalence is decidable for deterministic processes with no else branches and for the family of convergent subterm equational theories [Cortier and Delaune 2009]. Convergent subterm theories capture most standard primitives including asymmetric and symmetric encryption, hashes, signatures, and macs. A simplified proof of Baudet [2005] has been proposed by Chevalier and Rusinowitch [2012].
- Tiu and Dawson [2010] have designed and implemented a procedure for open bisimulation, a notion of equivalence stronger than the standard notion of trace equivalence. This procedure only works for a limited class of processes without else branches and for symmetric encryption and pairs only.
- Cheval et al. [2011] have proposed and implemented a procedure for trace equivalence and for a quite general class of processes that use standard primitives (symmetric and asymmetric encryption, hashes, signatures, pairs). In particular, this is the only decidability result that can consider nondeterministic processes and else branches.

However, these decidability results analyze equivalence for a *bounded number of sessions* only, that is, assuming that protocols are executed a limited number of times. This is of course a strong limitation. Even if no flaw is found when a protocol is executed n times, there is absolutely no guarantee that the protocol remains secure when it is executed $n + 1$ times. And actually, the existing tools for a bounded number of sessions can only analyze protocols for a very limited number of sessions, typically two or three. Another approach consists of implementing a procedure that is not guaranteed to terminate. This is in particular the case of ProVerif [Blanchet 2001], a well-established tool for checking security of protocols. ProVerif is able to check equivalence, although it does not always succeed [Blanchet et al. 2005]. It can check equivalence of *biprocesses*, that is, of two processes that have the same structure. ProVerif has been recently extended [Cheval and Blanchet 2013] to handle more processes, in particular with else branches, but it still cannot consider processes with very different structures. Of course, ProVerif does not correspond to any decidability result.

Our contribution. We study the decidability of equivalence of security protocols for an unbounded number of sessions. Even in the case of reachability properties such as secrecy, the problem is undecidable in general. In the past, several decidable fragments and semidecision procedures have been proposed for secrecy and authentication for an unbounded number of sessions. Our goal is to obtain analogous results in the case of equivalence properties. We therefore focus on a class of protocols for which secrecy

is decidable [Comon-Lundh and Cortier 2003]. This class, called *ping-pong protocols*, typically assumes that each protocol rule manipulates at most one variable and that the protocol is formed from a set of independent in/out rules. Intuitively, this corresponds to the assumption that, at each step of the protocol, upon receiving a message, there is at most one part of it that is unknown to the agent (typically a key, a nonce, or an encrypted packet).

Surprisingly, while this class is decidable for reachability, even a fragment of it (with only symmetric encryption) turns out to be undecidable for equivalence properties. We consequently further assume our protocols to be deterministic (i.e., given an input, there is at most one possible output). We show that equivalence is decidable for an unbounded number of sessions and for protocols with randomized symmetric and asymmetric encryption and with signatures. Since we need to assume that our constructors are randomized and since we assume “at most one variable,” we can only handle a very limited notion of (randomized) concatenation that appends atomic values.

Interestingly, we show that checking for equivalence of protocols actually amounts to checking equality of languages of deterministic pushdown automata. The language equivalence problem for deterministic pushdown automata is a difficult problem, shown to be decidable at ICALP in 1997 [Sénizergues 1997]. We actually characterize equivalence of protocols in terms of equivalence of deterministic generalized real-time pushdown automata, that is, deterministic pushdown automata with no epsilon-transition but such that the automata may unstack several symbols at a time. More precisely, we show how to associate to a process P an automata \mathcal{A}_P such that two processes are equivalent if and only if their corresponding automata yield the same language, and reciprocally, we show how to associate to an automata \mathcal{A} a process $P_{\mathcal{A}}$ such that two automata yield the same language if and only if their corresponding processes are equivalent, that is:

$$P \approx Q \Leftrightarrow L(\mathcal{A}_P) = L(\mathcal{A}_Q) \quad \text{and} \quad L(\mathcal{A}) = L(\mathcal{B}) \Leftrightarrow P_{\mathcal{A}} \approx P_{\mathcal{B}}.$$

Therefore, checking for equivalence of protocols is as difficult as checking for equivalence of deterministic generalized real-time pushdown automata.

To transform equivalence of processes into equivalence of pushdown automata, we first show how to get rid of an active attacker. More precisely, we show that

$$P \approx Q \Leftrightarrow P' \approx_{\text{fwd}} Q',$$

where \approx_{fwd} intuitively represents equivalence of processes when the attacker may only *forward* messages. This equivalence is obtained by partially encoding the attacker in P' and Q' , still preserving equivalence.

The decision procedure for checking equivalence of deterministic pushdown automata has been recently implemented by Henry and Sénizergues [2013]. We have therefore implemented our transformation from processes to pushdown automata, yielding the first tool that decides equivalence of (some class of) protocols for an unbounded number of sessions. As an application, we have analyzed several protocols of the literature, including a simplified version of the basic access control protocol (BAC) of the biometric passport [ICAO 2008].

We introduce the process algebra and its semantics in Section 2. We characterize the notion of ping-pong protocols and state our main results in Section 3. Sections 4 and 5 are devoted to decidability. More precisely, we show in Section 4 how to get rid of an active attacker by encoding it directly in the process. Next, we show in Section 5 how to encode equivalence between processes (in the presence of a forwarder attacker) into equivalence of pushdown automata, characterizing further which cases may result in nonequivalence. Finally, we study in Section 6 the converse translation and show that equivalence of pushdown automata can be reduced to equivalence of protocols. We

present our implementation and its application to protocols in Section 7. Concluding remarks can be found in Section 8.

2. MODEL FOR SECURITY PROTOCOLS

Security protocols are modeled through a process algebra that manipulates terms. We first give the syntax of our calculus in Section 2.1, before describing its semantics in Section 2.2. Then, in Section 2.3, we define the notion of equivalence of processes.

2.1. Syntax

Term Algebra. As usual, messages are represented by terms. More specifically, we consider a *sorted signature* with six sorts, `rand`, `key`, `msg`, `SymKey`, `PrivKey`, and `PubKey`, that represent, respectively, random numbers, keys, messages, symmetric keys, private keys, and public keys. We assume that `msg` subsumes the five other sorts, and `key` subsumes `SymKey`, `PrivKey`, and `PubKey`. We consider six function symbols, `senc` and `sdec`, `aenc` and `adec`, and `sign` and `check`, that represent symmetric, asymmetric encryption, and decryption as well as signatures. Since we are interested in the analysis of indistinguishability properties, we consider a randomized encryption scheme:

$$\begin{array}{ll} \text{senc} : \text{msg} \times \text{SymKey} \times \text{rand} \rightarrow \text{msg} & \text{sdec} : \text{msg} \times \text{SymKey} \rightarrow \text{msg} \\ \text{aenc} : \text{msg} \times \text{PubKey} \times \text{rand} \rightarrow \text{msg} & \text{adec} : \text{msg} \times \text{PrivKey} \rightarrow \text{msg} \\ \text{sign} : \text{msg} \times \text{PrivKey} \times \text{rand} \rightarrow \text{msg} & \text{check} : \text{msg} \times \text{PubKey} \rightarrow \text{msg}. \end{array}$$

We discuss in Section 7 how we can handle a limited notion of (randomized) concatenation.

We further assume an infinite set Σ_0 of *constant symbols* of sort `key` or `msg`, an infinite set Ch of constant symbols of sort `channel`, two infinite sets of *variables* \mathcal{X} , \mathcal{W} , and an infinite set of *names* $\mathcal{N} = \mathcal{N}_{\text{pub}} \uplus \mathcal{N}_{\text{prv}}$ of sort `rand`: \mathcal{N}_{pub} represents the random numbers drawn by the attacker, while \mathcal{N}_{prv} represents the random numbers drawn by the protocol's participants.

As usual, *terms* are defined as names, variables, and function symbols applied to other terms. We denote by $\mathcal{T}(\mathcal{F}, \mathcal{N}, \mathcal{X})$ the set of terms built on function symbols in \mathcal{F} , names in \mathcal{N} , and variables in \mathcal{X} . We simply write $\mathcal{T}(\mathcal{F}, \mathcal{N})$ when $\mathcal{X} = \emptyset$. We consider three particular signatures:

$$\begin{aligned} \Sigma_{\text{pub}} &= \{\text{senc}, \text{sdec}, \text{aenc}, \text{adec}, \text{sign}, \text{check}, \text{start}\} \\ \Sigma^+ &= \Sigma_{\text{pub}} \cup \Sigma_0 \quad \Sigma = \{\text{senc}, \text{aenc}, \text{sign}, \text{start}\} \cup \Sigma_0, \end{aligned}$$

where `start` $\notin \Sigma_0$ is a constant symbol of sort `msg`. The signature Σ_{pub} represents the functions/data available to the attacker, including a constant `start` used to start sessions of the protocols. The signature Σ^+ is the most general signature, while Σ models actual messages (with no failed computation). We assume a bijection between elements of sort `PrivKey` and `PubKey`. If k is a constant of sort `PrivKey`, k^{-1} will denote its image by this function, called *inverse*. The inverse of the inverse function is also denoted by $^{-1}$, so that $(k^{-1})^{-1} = k$. To keep homogeneous notations, we extend this function to symmetric keys: if k is of sort `SymKey`, then $k^{-1} = k$. The relation between encryption and decryption is represented through the following rewriting rules, yielding a convergent rewrite system:

$$\begin{aligned} \text{sdec}(\text{senc}(x, k_1, z), k_1) &\rightarrow x \\ \text{adec}(\text{aenc}(x, k_2, z), k_2^{-1}) &\rightarrow x \quad \text{check}(\text{sign}(x, k_3, z), k_3^{-1}) \rightarrow x, \end{aligned}$$

with k_1 of sort `SymKey`, k_2 of sort `PubKey`, and k_3 of sort `PrivKey`. For instance, the first rule models the fact that the decryption of a ciphertext will return the associated

plaintext when the right key is used to perform decryption. The two last rules are used to model asymmetric encryption and signatures. We denote by $t \downarrow$ the *normal form* of a term $t \in \mathcal{T}(\Sigma^+, \mathcal{N}, \mathcal{X})$.

Example 2.1. The term $m = \text{senc}(s, k, r)$ represents an encryption of the constant s with the key k using the random $r \in \mathcal{N}$, whereas $t = \text{sdec}(m, k)$ models the application of the decryption algorithm on m using k . We have that $t \downarrow = s$.

An attacker may build his or her own messages by applying functions to terms he or she already knows. Formally, a computation done by the attacker is modeled by a *recipe*, that is, a term in $\mathcal{T}(\Sigma_{\text{pub}}, \mathcal{N}_{\text{pub}}, \mathcal{W})$. The variables in \mathcal{W} intuitively refer to variables used to store messages learned by the attacker.

Process Algebra. The intended behavior of a protocol can be modeled by a *process* defined by the following grammar:

$P, Q := 0$	null process
$\text{in}(c, u).P$	input
$\text{out}(c, u).P$	output
$(P \mid Q)$	parallel
$!P$	replication
$\text{new } n.P$	name generation

where $u \in \mathcal{T}(\Sigma, \mathcal{N}, \mathcal{X})$, $n \in \mathcal{N}$, and $c \in Ch$.

The process 0 does nothing, and we sometimes omit it. The process “ $\text{in}(c, u).P$ ” expects a message m of the form u on channel c and then behaves like $P\theta$, where θ is a substitution such that $m = u\theta$. The process “ $\text{out}(c, u).P$ ” emits u on channel c and then behaves like P . The variables that occur in u are instantiated when the evaluation takes place. The process $P \mid Q$ runs P and Q in parallel. The process $!P$ executes P some arbitrary number of times. The process $\text{new } n.P$ invents a new name n and continues as P .

We write $fv(P)$ for the set of *free variables* that occur in P , that is, the set of variables that are not in the scope of an input. A *protocol* is a ground process, that is, a process P such that $fv(P) = \emptyset$.

Example 2.2. We consider a simplified version of the protocol presented in Denning and Sacco [1981]. The purpose of this protocol informally described next is to establish a key k_{AB} between two participants A and B using public key encryption and signature.

1. $A \rightarrow B : \text{aenc}(\text{sign}(k_{AB}, \text{sk}_A, r_A^1), \text{pk}_B, r_A^2)$
2. $B \rightarrow A : \text{ack}$

The agent A sends a symmetric key k_{AB} signed with A 's private key sk_A (using a fresh random number r_A^1), and the resulting ciphertext is encrypted with B 's public key pk_B (using a fresh random number r_A^2). The agent B answers this request by decrypting this message and verifying the signature. If all checks succeed, B informs the agent A by sending an acknowledgment, that is, the constant ack . The agents A and B can now use the symmetric key k_{AB} to communicate.

The role of A is modeled by a process P_A , while the role of B is modeled by P_B . We have that

$$P_A \stackrel{\text{def}}{=} !\text{in}(c_A, \text{start}).\text{new } r_A^1.\text{new } r_A^2.\text{out}(c_A, \text{aenc}(\text{sign}(k_{AB}, \text{sk}_A, r_A^1), \text{pk}_B, r_A^2)) \quad (1)$$

$$\mid !\text{in}(c'_A, \text{start}).\text{new } r_A^1.\text{new } r_A^2.\text{out}(c'_A, \text{aenc}(\text{sign}(k_{AC}, \text{sk}_A, r_A^1), \text{pk}_C, r_A^2)) \quad (2)$$

$$P_B \stackrel{\text{def}}{=} !\text{in}(c_B, \text{aenc}(\text{sign}(x, \text{sk}_A, z_1), \text{pk}_B, z_2)).\text{out}(c_B, \text{ack}). \quad (3)$$

The constants c_A, c'_A , and c_B are constants of sort channel; ack is a constant of sort msg ; and the constants $k_{AB}, k_{AC}, \text{sk}_A, \text{sk}_B, \text{sk}_C, \text{pk}_A, \text{pk}_B, \text{pk}_C$, which are in Σ_0 , are such that

- k_{AB}, k_{AC} are of sort SymKey ;
- $\text{sk}_A, \text{sk}_B, \text{sk}_C$ are of sort PrivKey ; and
- $\text{pk}_A, \text{pk}_B, \text{pk}_C$ of sort PubKey .

Moreover, we have that $\text{sk}_X^{-1} = \text{pk}_X$ for $X \in \{A, B, C\}$, whereas $k_{AB}^{-1} = k_{AB}$ and $k_{AC}^{-1} = k_{AC}$. Finally, r_A^1, r_A^2 are names of sort rand , and x (z_1, z_2 , respectively) is a variable of sort msg (rand , respectively).

Intuitively, P_A sends k_{AB} signed with sk_A and encrypted with pk_B to the agent B (branch 1). More generally, the agent A can start different sessions with different agents. Thus, the process P_A models the agent A initiating a session with B (branch 1) as well as with C (branch 2). The process P_B models the agent B answering a request from A . We could also consider the scenario where the agent B is also willing to talk to C or where the initiator, here played by A , is also played by other agents such as B . We consider here only a simpler case to keep the example reasonably short.

To model the whole protocol, we sent the public key $\text{pk}_A, \text{pk}_B, \text{pk}_C$ in clear, as well as the private key sk_C , to model the fact that the attacker may learn the private keys of some corrupted agents. This is modeled through the following process P_{key} :

$$P_{\text{key}} \stackrel{\text{def}}{=} !\text{in}(c_1, \text{start}).\text{out}(c, \text{pk}_A) \mid !\text{in}(c_2, \text{start}).\text{out}(c, \text{pk}_B) \mid \\ !\text{in}(c_3, \text{start}).\text{out}(c, \text{pk}_C) \mid !\text{in}(c_4, \text{start}).\text{out}(c, \text{sk}_C).$$

Then, the whole protocol is given by P , where P_A, P_B , and P_{key} evolve in parallel:

$$P \stackrel{\text{def}}{=} P_A \mid P_B \mid P_{\text{key}}.$$

This protocol is actually insecure as demonstrated by the following attack:

1. $A \rightarrow C$: $\text{aenc}(\text{sign}(k_{AC}, \text{sk}_A, r_A^1), \text{pk}_C, r_A^2)$
2. $C(A) \rightarrow B$: $\text{aenc}(\text{sign}(k_{AC}, \text{sk}_A, r_A^1), \text{pk}_B, r_C^1)$
3. $B \rightarrow A$: ack .

A initiates a session with a malicious user C , sending him or her a key k_{AC} . This malicious user then legally learns k_{AC} but also its signature $\text{sign}(k_{AC}, \text{sk}_A, r_A^1)$ under the signing key of A . He or she may then resend this key to B in the name of A . The agent B accepts the key k_{AC} as being a secret key between A and B .

2.2. Semantics

A *configuration* of a protocol is a pair $(P; \sigma)$, where

- \mathcal{P} is a multiset of processes. We often write $P \cup \mathcal{P}$, or $P \mid \mathcal{P}$, instead of $\{P\} \cup \mathcal{P}$.
- $\sigma = \{w_1 \triangleright m_1, \dots, w_n \triangleright m_n\}$ is a *frame*, that is, a substitution where w_1, \dots, w_n are variables in \mathcal{W} , and m_1, \dots, m_n are terms in $\mathcal{T}(\Sigma, \mathcal{N})$. Those terms represent the messages that are known by the attacker.

The operational semantics of protocol is defined by the relation $\xrightarrow{\alpha}$ over configurations described in Figure 1. For the sake of simplicity, we often write P instead of $(P; \emptyset)$.

The first rule (IN) allows the attacker to make a process progress by feeding it with a term he or she built with publicly available terms and symbols. The second one (OUT) lets the attacker gain knowledge of a message as soon as it is sent by a process: the corresponding message is added to the substitution of the current configuration. These

$$\begin{aligned}
& (\text{in}(c, u).P \cup \mathcal{P}; \sigma) \xrightarrow{\text{in}(c, R)} (P\theta \cup \mathcal{P}; \sigma) \quad (\text{IN}) \\
& \quad \text{where } R \text{ is a recipe such that } R\sigma \downarrow \in \mathcal{T}(\Sigma, \mathcal{N}) \text{ and } R\sigma \downarrow = u\theta \text{ for some } \theta \\
& (\text{out}(c, u).P \cup \mathcal{P}; \sigma) \xrightarrow{\text{out}(c, w_{i+1})} (P \cup \mathcal{P}; \sigma \cup \{w_{i+1} \triangleright u\}), \quad (\text{OUT}) \\
& \quad \text{where } i \text{ is the number of elements in } \sigma \\
& (!P \cup \mathcal{P}; \sigma) \xrightarrow{\tau} (P \cup !P \cup \mathcal{P}; \sigma) \quad (\text{REPL}) \\
& (\text{new } n.P \cup \mathcal{P}; \sigma) \xrightarrow{\tau} (P\{n'/n\} \cup \mathcal{P}; \sigma), \quad (\text{NEW}) \\
& \quad \text{where } n' \text{ is a fresh name in } \mathcal{N}_{\text{prv}}
\end{aligned}$$

Fig. 1. Operational semantics.

two rules are the only observable actions. The two remaining rules are quite standard and are unobservable (τ action) from the point of view of the attacker.

The relation $\xrightarrow{\text{tr}}$ between configurations (where tr is a sequence of actions) is defined in a usual way as the reflexive and transitive closure of the relation $\xrightarrow{\alpha}$. Given a sequence of observable actions tr , we write $K \xrightarrow{\text{tr}} K'$ when there exists tr' such that $K \xrightarrow{\text{tr}'} K'$ and w is obtained from tr' by erasing all occurrences of τ . For every configuration K , we define its *set of traces* as follows:

$$\text{trace}(K) = \{(\text{tr}, \sigma) \mid K \xrightarrow{\text{tr}} (P; \sigma) \text{ for some configuration } (P; \sigma)\}.$$

Example 2.3. Going back to the protocol introduced in Example 2.2, we consider the scenario corresponding to the attack.

- (1) The public keys of all the participants are disclosed as well as the secret key sk_C of the corrupted agent C . Formally, let $K_0 \stackrel{\text{def}}{=} (P; \emptyset)$; we have that

$$K_0 \xrightarrow{\text{in}(c_1, \text{start}).\text{out}(c_1, w_1).\text{in}(c_2, \text{start}).\text{out}(c_2, w_2).\text{in}(c_3, \text{start}).\text{out}(c_3, w_3)} \xrightarrow{\text{in}(c_4, \text{start}).\text{out}(c_4, w_4)} (P; \sigma_0),$$

where $\sigma_0 = \{w_1 \triangleright \text{pk}_A, w_2 \triangleright \text{pk}_B, w_3 \triangleright \text{pk}_C, w_4 \triangleright \text{sk}_C\}$.

- (2) The agent A initiates a session with C and sends the corresponding encrypted message. More formally, we have that

$$(P; \sigma_0) \xrightarrow{\text{in}(c'_A, \text{start}).\text{out}(c'_A, w_5)} (P; \sigma),$$

where $\sigma = \sigma_0 \cup \{w_5 \triangleright \text{aenc}(\text{sign}(\text{k}_{AC}, \text{sk}_A, r_A^1), \text{pk}_C, r_A^2)\}$ and r_A^1, r_A^2 are (fresh) names in \mathcal{N}_{prv} .

Hence, we have that $(\text{tr}, \sigma) \in \text{trace}(K_0)$, where

$$\begin{aligned} \text{tr} = & \text{in}(c_1, \text{start}).\text{out}(c_1, w_1).\text{in}(c_2, \text{start}).\text{out}(c_2, w_2).\text{in}(c_3, \text{start}).\text{out}(c_3, w_3). \\ & \text{in}(c_4, \text{start}).\text{out}(c_4, w_4).\text{in}(c'_A, \text{start}).\text{out}(c'_A, w_5). \end{aligned}$$

In this execution trace, first the keys $\text{pk}_A, \text{pk}_B, \text{pk}_C$, and sk_C are sent after having called the corresponding process. Then, branch (2) of P is triggered.

2.3. Trace Equivalence

Intuitively, two processes are equivalent if they cannot be distinguished by any attacker. Trace equivalence can be used to formalize many interesting security properties, in particular privacy-type properties, such as those studied, for instance, in Arapinis et al. [2010] and Bruso et al. [2010]. We first introduce a notion of the intruder's knowledge well suited to cryptographic primitives for which the success of decrypting or checking a signature is visible.

Definition 2.4. Two frames σ_1 and σ_2 are *statically equivalent*, $\sigma_1 \sim \sigma_2$, when we have that $\text{dom}(\sigma_1) = \text{dom}(\sigma_2)$, and

- for any recipe R , $R\sigma_1 \downarrow \in \mathcal{T}(\Sigma, \mathcal{N})$ if and only if $R\sigma_2 \downarrow \in \mathcal{T}(\Sigma, \mathcal{N})$; and
- for all recipes R_1 and R_2 such that $R_1\sigma_1 \downarrow, R_2\sigma_1 \downarrow \in \mathcal{T}(\Sigma, \mathcal{N})$, we have that $R_1\sigma_1 \downarrow = R_2\sigma_1 \downarrow$ if and only if $R_1\sigma_2 \downarrow = R_2\sigma_2 \downarrow$.

Intuitively, two frames are equivalent if an attacker cannot see the difference between the two situations they represent: if some computation fails in σ_1 , it should fail in σ_2 as well, and σ_1 and σ_2 should satisfy the same equalities.

Example 2.5. Consider the two following frames:

- (1) $\sigma_1 \stackrel{\text{def}}{=} \sigma_0 \cup \{w_5 \triangleright \text{aenc}(\text{sign}(k_{AC}, \text{sk}_A, r_A^1), \text{pk}_C, r_A^2), w_6 \triangleright k_{AC}\},$
- (2) $\sigma_2 \stackrel{\text{def}}{=} \sigma_0 \cup \{w_5 \triangleright \text{aenc}(\text{sign}(k_{AC}, \text{sk}_A, r_A^1), \text{pk}_C, r_A^2), w_6 \triangleright k\},$

where k is a (private) constant in Σ_0 . We have that $\sigma_1 \not\sim \sigma_2$. Indeed, consider the recipes $R_1 = \text{check}(\text{adec}(w_5, w_4), w_1)$ and $R_2 = w_6$. We have that $R_1\sigma_1 \downarrow = R_2\sigma_1 \downarrow = k_{AC}$, whereas $R_1\sigma_2 \downarrow = k_{AC}$ and $R_2\sigma_2 \downarrow = k$ and thus $R_1\sigma_2 \downarrow \neq R_2\sigma_2 \downarrow$.

Intuitively, two processes are *trace equivalent* if, however they behave, the resulting sequences of messages observed by the attacker are in static equivalence.

Definition 2.6. Let P and Q be two protocols. We have that $P \sqsubseteq Q$ if for every $(\text{tr}, \sigma) \in \text{trace}(P)$, there exists $(\text{tr}', \sigma') \in \text{trace}(Q)$ such that $\text{tr} = \text{tr}'$ and $\sigma \sim \sigma'$. They are *trace equivalent*, written $P \approx Q$, if $P \sqsubseteq Q$ and $Q \sqsubseteq P$.

Example 2.7. Continuing Example 2.2, our naive protocol is secure if the key received by B remains private. To model this, we modify the process P_B as follows:

$$\begin{aligned} P_B^l &\stackrel{\text{def}}{=} !\text{in}(c_B, \text{aenc}(\text{sign}(x, \text{sk}_A, z_1), \text{pk}_B, z_2)).\text{out}(c_B, x) \\ P_B^r &\stackrel{\text{def}}{=} !\text{in}(c_B, \text{aenc}(\text{sign}(x, \text{sk}_A, z_1), \text{pk}_B, z_2)).\text{out}(c_B, k). \end{aligned}$$

Then, to model secrecy of the key received by B , we consider the following equivalence: $P_A \mid P_B^l \mid P_{\text{key}} \approx P_A \mid P_B^r \mid P_{\text{key}}$. An attacker should not distinguish between two instances of the protocol, one where B used the key established through the protocol and one where a magic key k is used instead.

However, our protocol is insecure. An attacker may easily learn k_{AC} and send to B a message of the expected form (as if it were issued by A) that will contain this corrupted key instead of k_{AB} . Formally, we have that

$$P_A \mid P_B^l \mid P_{\text{key}} \not\approx P_A \mid P_B^r \mid P_{\text{key}}.$$

This is reflected by the trace tr' described as follows:

$$\text{tr}' \stackrel{\text{def}}{=} \text{tr}.\text{in}(c_B, \text{aenc}(\text{adec}(w_5, \text{sk}_C), w_2, r_C)).\text{out}(c_B, w_6),$$

where r_C is a name in \mathcal{N}_{pub} .

We have that $(\text{tr}', \sigma_1) \in \text{trace}(K_0)$ with $K_0 = (P_A \mid P_B^l \mid P_{\text{key}}; \sigma_1)$ and σ_1 as defined in Example 2.5. Because of the existence of only one branch using each channel, there is only one possible execution of $P_A \mid P_B^r \mid P_{\text{key}}$ (up to a bijective renaming of the private names of sort rand) matching the labels in tr' , and the corresponding execution will allow us to reach the frame σ_2 as described in Example 2.5. We have already seen that static equivalence does not hold, that is, $\sigma_1 \not\sim \sigma_2$.

3. PING-PONG PROTOCOLS

We aim at providing a decidability result for the problem of trace equivalence between protocols in the presence of replication. However, it is well known that replication leads to undecidability even for the simple case of reachability properties. Thus, we consider a class of protocols, called C_{pp} , for which (in a slightly different setting), reachability has already been proved decidable [Comon-Lundh and Cortier 2003].

3.1. Class C_{pp}

We basically consider ping-pong protocols (an output is computed using only the message previously received in input), and we assume a kind of determinism. Moreover, we restrict the terms that are manipulated throughout the protocols: only one unknown message (modeled by the use of a variable of sort `msg`) can be received at each step.

We fix a variable $x \in \mathcal{X}$ of sort `msg`. An *input term* (*output term*, respectively) is a term defined by the following grammars given:

$$u := x \mid s \mid f(u, k, z) \quad v := x \mid s \mid f(v, k, r),$$

where $s, k \in \Sigma_0 \cup \{\text{start}\}$, $z \in \mathcal{X}$, $f \in \{\text{senc}, \text{aenc}, \text{sign}\}$, and $r \in \mathcal{N}$. Intuitively, no destructor should be used explicitly. Moreover, we assume that each variable (name, respectively) occurs at most once in u (v , respectively).

Definition 3.1. C_{pp} is the class of protocol of the form

$$P = \prod_{i=1}^n \prod_{j=1}^{p_i} \text{in}(c_i, u_i^j). \text{new } r_1. \dots \text{new } r_{k_j}. \text{out}(c_i, v_i^j)$$

such that

- (1) for all $i \in \{1, \dots, n\}$ and $j \in \{1, \dots, p_i\}$, $k_j^j \in \mathbb{N}$, u_i^j is an input term, and v_i^j is an output term where names occurring in v_i^j are included in $\{r_1, \dots, r_{k_j^j}\}$;
- (2) for all $i \in \{1, \dots, n\}$ and $j_1, j_2 \in \{1, \dots, p_i\}$, if $j_1 \neq j_2$, then for any renaming of variables, $u_i^{j_1}$ and $u_i^{j_2}$ are not unifiable.¹

Each subprocess $\text{in}(c_i, u_i^j). \text{new } r_1. \dots \text{new } r_{k_j^j}. \text{out}(c_i, v_i^j)$ is called a *branch* of P .

Item 1 holds for any process representing a protocol: the variables of the output should be bound by the input. Item 2 enforces a deterministic behavior: a particular input action can only be accepted by one branch of the protocol. This is a natural restriction since most of the protocols are indeed deterministic: an agent should usually know exactly what to do once he or she has received a message. Actually, the main limitations of the class C_{pp} is that we consider a restricted signature (e.g., no pair, no hash function), and names can only be used to produce randomized ciphertexts.

Example 3.2. The protocols described in Example 2.7 are in C_{pp} . For instance, we can check that

- $\text{aenc}(\text{sign}(x, \text{sk}_A, z_1), \text{pk}_B, z_2)$ is an input term, and
- $\text{aenc}(\text{sign}(k_{AB}, \text{sk}_A, r_A^1), \text{pk}_B, r_B^2)$ is an output term.

Moreover, the determinism condition (item 2) is clearly satisfied: each branch of the protocol $P_A \mid P_B^l \mid P_{\text{key}}$ ($P_A \mid P_B^r \mid P_{\text{key}}$, respectively) uses a different channel.

When studying trace equivalence (or even trace inclusion), we can even safely force a process to perform an input action followed directly by its associated output action.

¹That is, there does not exist θ such that $u_i^{j_1}\theta = u_i^{j_2}\theta$.

We consider a set of “big-step” traces, defined as follows:

$$\text{trace}^{\text{io}*}(K) = \left\{ (\text{tr}, \sigma) \mid K \xRightarrow{\text{tr}} (\mathcal{P}; \sigma) \text{ for some configuration } (\mathcal{P}; \sigma) \right. \\ \left. \text{with tr sequence of input-output blocks.} \right\}$$

The notion of trace inclusion (trace equivalence, respectively) w.r.t. big-step traces is defined accordingly.

Definition 3.3. Let P and Q be two protocols. We have that $P \sqsubseteq^{\text{io}*} Q$ if for every $(\text{tr}, \sigma) \in \text{trace}^{\text{io}*}(P)$, there exists $(\text{tr}', \sigma') \in \text{trace}^{\text{io}*}(Q)$ such that $\text{tr} = \text{tr}'$ and $\sigma \sim \sigma'$. They are *trace equivalent*, written $P \approx^{\text{io}*} Q$, if $P \sqsubseteq^{\text{io}*} Q$ and $Q \sqsubseteq^{\text{io}*} P$.

Due to the form of protocols in \mathcal{C}_{pp} , any trace made up of input and output actions can first be completed with all the available output actions and then be mapped to a trace that is made up of input-output blocks only. Thus, we have that the two notions of trace equivalence coincide.

PROPOSITION 3.4. Let P and Q be two protocols in \mathcal{C}_{pp} . We have that $P \sqsubseteq^{\text{io}*} Q$ if and only if $P \sqsubseteq Q$.

This proposition easily follows from the fact that for any process of \mathcal{C}_{pp} , any input is immediately followed by an output.

3.2. Main Results

Our first main contribution is a decision procedure for trace equivalence of processes in \mathcal{C}_{pp} .

THEOREM 3.5. Let P and Q be two protocols in \mathcal{C}_{pp} . The problem whether P and Q are trace equivalent, that is, $P \approx Q$, is decidable.

Deciding trace equivalence is done in two main steps:

- (1) First, we show how to reduce trace equivalence between protocols in \mathcal{C}_{pp} to the problem of deciding trace equivalence (still between protocols in \mathcal{C}_{pp}) when the attacker acts as a *forwarder*, that is, when the attacker may only forward messages obtained through the protocol. This step is detailed in Section 4.
- (2) Then, we encode the problem of deciding trace equivalence for forwarding attackers into the problem of language equivalence for real-time generalized pushdown deterministic automata (GPDA), that is, deterministic pushdown automata with no epsilon-transition but such that the automata may unstack several symbols at a time. This step is detailed in Section 5.

We also provide an implementation of our translation from protocols to pushdown automata, yielding a tool for automatically checking equivalence of security protocols for an unbounded number of sessions. This contribution is described in Section 7.

Actually, we characterize equivalence of protocols in terms of equivalence of GPDA. Indeed, Step (2) shows how to associate to a process P an automata \mathcal{A}_P such that two processes are equivalent if and only if their corresponding automata yield the same language. Conversely, we also show how to associate to an automata \mathcal{A} a process $P_{\mathcal{A}}$ such that two automata yield the same language if and only if their corresponding processes are equivalent. This reverse encoding from pushdown automata to protocols is explained in Section 6.

Our second contribution is an undecidability result. The class \mathcal{C}_{pp} is somewhat limited, but extending \mathcal{C}_{pp} to nondeterministic processes immediately yields undecidability of trace equivalence. More precisely, we have that trace inclusion of processes in \mathcal{C}_{pp} is undecidable.

THEOREM 3.6. *The following problem is undecidable:*

Input. P and Q two protocols in C_{pp} .

Output. Whether P is trace included in Q , that is, $P \sqsubseteq Q$.

A direct encoding of the Post Correspondence Problem (PCP) into an inclusion of two protocols of this class is given in Appendix A. Alternatively, this undecidability result is also a consequence of the reduction result established in Section 6 and the undecidability result established in Friedman [1976]. Nonetheless, we present in Appendix A the direct encoding of PCP into protocol equivalence since some ideas might be reused to show undecidability of trace equivalence for some other classes, whereas the alternative proof required a first encoding to transform a protocol into a pushdown automaton.

Undecidability of trace inclusion actually implies undecidability of trace equivalence as soon as processes are nondeterministic. Indeed, consider the choice operator $+$ whose (standard) semantics is given by the following rules:

$$(\{P + Q\} \cup P; \sigma) \xrightarrow{\tau} (P \cup P; \sigma) \quad (\{P + Q\} \cup P; \sigma) \xrightarrow{\tau} (Q \cup P; \sigma).$$

COROLLARY 3.7. *Let P , Q_1 , and Q_2 be three protocols in C_{pp} . The problem whether P is equivalent to $Q_1 + Q_2$, that is, $P \approx Q_1 + Q_2$, is undecidable.*

Indeed, consider P and Q_1 , for which trace inclusion encodes PCP, and let $Q_2 = P$. Trivially, $P \sqsubseteq Q_1 + Q_2$. Thus, $P \approx Q_1 + Q_2$ if and only if $Q_1 + Q_2 \sqsubseteq P$, that is, if and only if $Q_1 \sqsubseteq P$, hence the undecidability result.

4. GETTING RID OF THE FULL ATTACKER

We show in this section how to reduce trace equivalence between protocols in C_{pp} to the problem of deciding trace equivalence (still between protocols in C_{pp}) when the attacker acts as a *forwarder*, that is, when the attacker may only forward messages obtained through the protocols. This new semantics induced a new notion of trace equivalence, denoted \approx_{fwd} , which is formally defined in Section 4.1.

To counterbalance the effects of this simple forwarder semantics, the key idea consists of modifying the protocols under study by adding new rules that encrypt and decrypt messages on demand for the forwarder. Formally, we define a transformation T_{fwd} (see Section 4.2) that associates to a pair of protocols in C_{pp} a finite set of pairs of protocols (still in C_{pp}), and we show the following result:

PROPOSITION 4.1. *Let P and Q be two protocols in C_{pp} . We have that*

$$P \approx Q \text{ if and only if } P' \approx_{\text{fwd}} Q' \text{ for some } (P', Q') \in T_{\text{fwd}}(P, Q).$$

4.1. Forwarder Semantics

We first define the actions of a forwarder by modifying our semantics. Roughly, we restrict the recipes R , R_1 , and R_2 that are used in the IN rule and in static equivalence (Definition 2.4) to be either the public constant start or a variable in \mathcal{W} . Intuitively, this corresponds to the fact that the forwarder attacker should no longer build a message on his or her own. This leads us to consider a new relation \rightarrow_{fwd} between configurations, which is the relation induced by the rules described in Figure 2.

The relations $\xrightarrow{\text{tr}}_{\text{fwd}}$ and $\xRightarrow{\text{tr}}_{\text{fwd}}$ between configurations where tr is a sequence of actions (observable actions, respectively) are defined as expected. For every configuration K ,

$$\begin{aligned}
& (\text{in}(c, u).P \cup \mathcal{P}; \sigma) \xrightarrow{\text{in}(c, R)}_{\text{fwd}} (P\theta \cup \mathcal{P}; \sigma), \\
& \quad \text{where } R \in \{\text{start}\} \cup \mathcal{W} \text{ and } R\sigma \downarrow = u\theta \text{ for some } \theta \\
& (\text{out}(c, u).P \cup \mathcal{P}; \sigma) \xrightarrow{\text{out}(c, w_{i+1})}_{\text{fwd}} (P \cup \mathcal{P}; \sigma \cup \{w_{i+1} \triangleright u\}), \\
& \quad \text{where } i \text{ is the number of elements in } \sigma \\
& (!P \cup \mathcal{P}; \sigma) \xrightarrow{\tau}_{\text{fwd}} (P \cup !P \cup \mathcal{P}; \sigma) \\
& (\text{new } n.P \cup \mathcal{P}; \sigma) \xrightarrow{\tau}_{\text{fwd}} (P\{n'/n\} \cup \mathcal{P}; \sigma) \quad \text{where } n' \text{ is a fresh name in } \mathcal{N}_{\text{prv}}
\end{aligned}$$

Fig. 2. Semantics for a forwarder attacker.

we define its *set of traces w.r.t. the forwarder semantics* as follows:

$$\text{trace}_{\text{fwd}}(K) = \left\{ (\text{tr}, \sigma) \mid \begin{array}{l} K \xrightarrow[\text{fwd}]{\text{tr}} (\mathcal{P}; \sigma) \text{ for some configuration } (\mathcal{P}; \sigma) \\ \text{with tr sequence of input-output blocks.} \end{array} \right\}$$

We also need to adapt our notion of static equivalence.

Definition 4.2. Two frames σ_1 and σ_2 are *statically equivalent w.r.t. the forwarder semantics*, denoted $\sigma_1 \sim_{\text{fwd}} \sigma_2$, when we have that $\text{dom}(\sigma_1) = \text{dom}(\sigma_2)$, and for all recipes R_1 and R_2 in $\{\text{start}\} \cup \mathcal{W}$, we have that $R_1\sigma_1 = R_2\sigma_1$ if and only if $R_1\sigma_2 = R_2\sigma_2$.

This induces a new notion of trace equivalence, which is formally defined as follows:

Definition 4.3. Let P and Q be two protocols. We have that $P \sqsubseteq_{\text{fwd}} Q$ if for every $(\text{tr}, \sigma) \in \text{trace}_{\text{fwd}}(P)$, there exists $(\text{tr}', \sigma') \in \text{trace}_{\text{fwd}}(Q)$ such that $\text{tr} = \text{tr}'$ and $\sigma \sim_{\text{fwd}} \sigma'$. They are *trace equivalent w.r.t. the forwarder semantics*, written $P \approx_{\text{fwd}} Q$, if $P \sqsubseteq_{\text{fwd}} Q$ and $Q \sqsubseteq_{\text{fwd}} P$.

Example 4.4. The trace exhibited in Example 2.3 is still a valid one according to the forwarder semantics, but the frames σ_1 and σ_2 described in Example 2.5 are now in equivalence w.r.t. \sim_{fwd} . Actually, we have that $P_A \mid P_B^l \mid P_{\text{key}} \approx_{\text{fwd}} P_A \mid P_B^r \mid P_{\text{key}}$. Indeed, the fact that a forwarder simply acts as a relay prevents him or her from mounting the aforementioned attack.

4.2. Toward a Forwarder Attacker

As illustrated in Example 4.4, the forwarder semantics is very restrictive: a forwarder cannot rely on his or her deduction capabilities to mount an attack. We show, however, that we can still restrict ourselves to trace equivalence w.r.t. a forwarder.

Intuitively, we transform any two processes P, Q into processes \tilde{P}, \tilde{Q} such that $P \approx Q$ if and only if $\tilde{P} \approx_{\text{fwd}} \tilde{Q}$. Roughly, this transformation consists of two steps:

- (1) First, we guess among the keys of the protocols P and the keys of the protocols Q those that are deducible by the attacker, as well as a bijection α between these two sets. We can show that such a bijection necessarily exists when $P \approx Q$.
- (2) Then, to compensate for the fact that the attacker is a simple forwarder, we give him or her access to encryption/decryption oracles for any deducible key k , adding branches in the processes.

To maintain the equivalence, we do a similar transformation in both P and Q relying on the bijection α . We ensure that the set of deducible keys has been correctly guessed by adding of some extra processes. Then the main step of the proof consists of showing that the forwarder has now the same power as a full attacker, even though he or she cannot reuse the same randomness in two distinct encryptions, as a real attacker could.

Example 4.5. To better illustrate this section, we consider a variant of the processes introduced in Section 2, where agent A is now willing to talk only to B :

$$P \stackrel{\text{def}}{=} P'_A \mid P'_B \mid P_{\text{key}} \quad Q \stackrel{\text{def}}{=} P'_A \mid P''_B \mid P_{\text{key}},$$

where P'_B , P''_B are defined in Example 2.7 and P_{key} is defined in Example 2.2, whereas P'_A is defined as follows (only the first branch of P_A):

$$P'_A \stackrel{\text{def}}{=} !\text{in}(c_A, \text{start}).\text{new } r_A^1.\text{new } r_A^2.\text{out}(c_A, \text{aenc}(\text{sign}(k_{AB}, \text{sk}_A, r_A^1), \text{pk}_B, r_A^2)). \quad (1)$$

This scenario excludes the aforementioned attack and we have that $P \approx Q$. This has been formally checked using our prototype (see Section 7).

4.2.1. Guessing Deducible Keys. The purpose of this section is to restrict our attention to protocols that explicitly disclose their deducible keys \mathcal{K}_P and \mathcal{K}_Q . Since we do not want to rely on a particular procedure for computing these two sets, the idea is to guess a possible superset of each set, namely, K and K' , and then ensure that these sets K and K' contain *at least* the deducible keys.

Definition 4.6. Let P be a protocol in \mathcal{C}_{pp} . A term t is *deducible* in P if there exists a trace $(\text{tr}, \phi) \in \text{trace}(P)$ and a recipe R (i.e., a term in $\mathcal{T}(\Sigma_{\text{pub}}, \mathcal{N}_{\text{pub}}, \mathcal{W})$) such that $R\phi \downarrow = t$.

Example 4.7. Continuing Example 4.5, we have that P and Q are in \mathcal{C}_{pp} . It is easy to notice that k_{AB} is deducible in P , whereas k is deducible in Q , since these keys are revealed at the end of B 's execution. For both P and Q , the trace $\text{tr} = \text{in}(c_A, \text{start}).\text{out}(c_A, w_1).\text{in}(c_B, w_1).\text{out}(c_B, w_2)$ and the recipe $R = w_2$ is a witness of this fact.

Two equivalent processes have the same set of deducible keys, up to some bijective renaming.

LEMMA 4.8. Let P and Q be two protocols in \mathcal{C}_{pp} and \mathcal{K}_P (\mathcal{K}_Q , respectively) be the set of deducible constants of sort *key* that occur in P (Q , respectively); if $P \approx Q$, then there exists a unique bijection α from \mathcal{K}_P to \mathcal{K}_Q such that for every trace $(\text{tr}, \phi) \in \text{trace}(P)$, there exists a trace $(\text{tr}, \psi) \in \text{trace}(Q)$ such that for any recipe R and any $k \in \mathcal{K}_P$:

— $R\phi \downarrow$ is of sort s if and only if $R\psi \downarrow$ is of sort s ,

where $s \in \{\text{SymKey}, \text{PubKey}, \text{PrivKey}\}$;

— $R\phi \downarrow = k$ if and only if $R\psi \downarrow = \alpha(k)$; and

— $R\phi \downarrow = k^{-1}$ if and only if $R\psi \downarrow = (\alpha(k))^{-1}$;

and conversely, for every $(\text{tr}, \psi) \in \text{trace}(Q)$, there exists a trace $(\text{tr}, \phi) \in \text{trace}(P)$ satisfying the same properties.

PROOF. (sketch) The relation α is defined as follows:

for every $k \in \mathcal{K}_P$ of sort s and every trace $(\text{tr}, \phi) \in \text{trace}(P)$ and recipe R such that $R\phi \downarrow = k$, we define $\alpha(k) = R\psi \downarrow$, where ψ is the only frame such that $(\text{tr}, \psi) \in \text{trace}(Q)$.

The existence of such a frame comes from the fact that $P \approx Q$, whereas its unicity is a consequence of the determinism of protocols in \mathcal{C}_{pp} .

Then, we show that this relation α is uniquely defined and satisfies all the requirements exploiting the strong relationship between P and Q through the relation $P \approx Q$. \square

Example 4.9. Continuing Example 4.5, we have $\mathcal{K}_P = \{\text{pk}_A, \text{pk}_B, \text{pk}_C, \text{sk}_C, k_{AB}\}$, whereas $\mathcal{K}_Q = \{\text{pk}_A, \text{pk}_B, \text{pk}_C, \text{sk}_C, k\}$. The unique bijection α mentioned in the previous lemma is defined as follows: $\alpha(k_{AB}) = k$, and $\alpha(k') = k'$ otherwise.

Definition 4.10. Let P be a protocol in \mathcal{C}_{pp} and K be a set of constants of sort key that occur in P . If for every $k \in K$ there exist a channel name c_k and a branch $!in(c_k, \text{start}).out(c_k, k)$ in P , then P is said to *disclose* K .

Example 4.11. Continuing our running example, P and Q clearly disclose $K = \{pk_A, pk_B, pk_C, sk_C\}$.

LEMMA 4.12. Let P and Q be two protocols in \mathcal{C}_{pp} and S (S' , respectively) the set of keys of P (Q , respectively). Then $P \approx Q$ if and only if there exist two sets $K \subseteq S$ and $K' \subseteq S'$ and a bijection $\alpha : K \rightarrow K'$ such that $\bar{P} \approx \bar{Q}$, where:

$$\begin{aligned} \bar{P} = P \mid & !in(c^0, \text{start}).out(c^0, 0) \mid !in(c^1, \text{start}).out(c^1, 1) \\ & \mid \mid_{k \in K} !in(c_{k, \alpha(k)}, \text{start}).out(c_{k, \alpha(k)}, k) \mid \mid_{k \in S \setminus K} !in(c, k).out(c, 0) \\ \bar{Q} = Q \mid & !in(c^0, \text{start}).out(c^0, 0) \mid !in(c^1, \text{start}).out(c^1, 1) \\ & \mid \mid_{k \in K} !in(c_{k, \alpha(k)}, \text{start}).out(c_{k, \alpha(k)}, \alpha(k)) \mid \mid_{k \in S' \setminus K'} !in(c, k).out(c, 1), \end{aligned}$$

and $0, 1$ are new constants and c^0, c^1 , the $c_{k, \alpha(k)}$, and c are fresh channels.

Moreover, assuming the existence of such sets and bijection such that $\bar{P} \approx \bar{Q}$, the two protocols are disclosing their deducible keys.

We call $\mathcal{T}_{key}(P, Q)$ the set of such pairs (\bar{P}, \bar{Q}) of modified protocols.

PROOF. Let \mathcal{K}_P (\mathcal{K}_Q , respectively) be the set of deducible constants of sort key that occur in P (Q , respectively). We prove the two directions separately.

(\Rightarrow) If $P \approx Q$, by Lemma 4.8, for $K = \mathcal{K}_P$ and $K' = \mathcal{K}_Q$, we get the existence of such a bijection α . Because keys in $S \setminus \mathcal{K}_P$ and $S' \setminus \mathcal{K}_Q$ are not deducible, the branches on channel c can never be triggered. Moreover, as $P \approx Q$, any trace of P (Q , respectively) inputting or outputting on a channel $c_{k, \alpha(k)}$ for k in \mathcal{K}_P can be matched in Q (P , respectively). Indeed, for every couple (k, k^{-1}) of deducible keys and for any recipe reducing to k (k^{-1} , respectively) in P , the same recipe reduces to $\alpha(k)$ ($\alpha(k)^{-1}$, respectively) in Q , thanks to the properties of α described in Lemma 4.8.

(\Leftarrow) For the converse implication, we first remark that necessarily we have that $\mathcal{K}_P \subseteq K$ and $\mathcal{K}_Q \subseteq K'$. Indeed, suppose there exists, for instance, $k \in \mathcal{K}_P \setminus K$. Since k is deducible, there exist a trace $(tr, \phi) \in \text{trace}(P)$ and a recipe R such that $R\phi \downarrow = k$. Since (tr, ϕ) is also a trace of \bar{P} , we consider the trace

$$tr' = tr.in(c, R).out(c, w_{|\phi|+1}).in(c^0, \text{start}).out(c^0, w_{|\phi|+2}).in(c^1, \text{start}).out(c^1, w_{|\phi|+3})$$

along with its frame $\phi' = \phi \cup \{w_{|\phi|+1} \triangleright 0, w_{|\phi|+2} \triangleright 0, w_{|\phi|+3} \triangleright 1\}$. If $\bar{P} \approx \bar{Q}$, then there exists $(tr', \psi') \in \text{trace}(\bar{Q})$ such that ϕ and ψ are statically equivalent. But any output on c in Q leads to the constant 1, breaking static equivalence. We conclude in a similar way in case $k \in \mathcal{K}_Q \setminus K'$.

Finally, we need to prove that $\bar{P} \approx \bar{Q}$ implies $P \approx Q$. For every trace $(tr, \phi) \in \text{trace}(P)$, $(tr, \phi) \in \text{trace}(\bar{P})$, and as $\bar{P} \approx \bar{Q}$, there exists a trace $(tr, \psi) \in \text{trace}(\bar{Q})$ such that ϕ is statically equivalent to ψ . Because c^0, c^1, c , and the $c_{k, \alpha(k)}$ are new channels, tr does not use transitions on those, and thus $(tr, \psi) \in \text{trace}(Q)$. The same goes for any trace of Q , hence showing the trace equivalence of P and Q . \square

Example 4.13. Continuing our example, let $K = \mathcal{K}_P$ and $K' = \mathcal{K}_Q$, and α is the bijection defined in Example 4.9. Checking equivalence of $P \approx Q$ amounts to checking

whether $\bar{P} \approx \bar{Q}$, where \bar{P} and \bar{Q} are defined as follows:

$$\begin{aligned}\bar{P} &= P \mid !\text{in}(c^0, \text{start}).\text{out}(c^0, 0) \mid !\text{in}(c^1, \text{start}).\text{out}(c^1, 1) \\ &\quad \mid !\text{in}(c_{k_{AB}, k}, \text{start}).\text{out}(c_{k_{AB}, k}, k_{AB}) \\ &\quad \mid !\text{in}(c, \text{sk}_A).\text{out}(c, 0) \mid !\text{in}(c, \text{sk}_B).\text{out}(c, 0) \\ \bar{Q} &= Q \mid !\text{in}(c^0, \text{start}).\text{out}(c^0, 0) \mid !\text{in}(c^1, \text{start}).\text{out}(c^1, 1) \\ &\quad \mid !\text{in}(c_{k_{AB}, k}, \text{start}).\text{out}(c_{k_{AB}, k}, k) \\ &\quad \mid !\text{in}(c, \text{sk}_A).\text{out}(c, 1) \mid !\text{in}(c, \text{sk}_B).\text{out}(c, 1).\end{aligned}$$

If $\bar{P} \approx \bar{Q}$, then sk_A and sk_B cannot be deducible and thus \bar{P} and \bar{Q} disclose their set of deducible keys.

4.2.2. Adding Oracles. To compensate for the fact that the attacker is a simple forwarder, we give him or her access to encryption/decryption oracles for any deducible key k , adding branches in the processes. We rely on the bijection α computed in the previous section to do this in a compatible way on both sides of the equivalence.

LEMMA 4.14. *Let P and Q be two protocols in \mathcal{C}_{pp} respectively disclosing two sets of keys K and K' as in Lemma 4.12. Then $P \approx Q$ if and only if $\bar{P} \approx_{\text{fwd}} \bar{Q}$, where*

$$\begin{aligned}\bar{P} &= P \mid \begin{array}{l} \mid_{k \in K^{\text{SymKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{senc}}, x).\text{new } n.\text{out}(c_{k, \alpha(k)}, \text{senc}(x, k, n)) \\ \mid_{k \in K^{\text{SymKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{sdec}}, \text{senc}(x, k, y)).\text{out}(c_{k, \alpha(k)}^{\text{sdec}}, x) \\ \mid_{k \in K^{\text{PubKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{aenc}}, x).\text{new } n.\text{out}(c_{k, \alpha(k)}^{\text{aenc}}, \text{aenc}(x, k, n)) \\ \mid_{k \in K^{\text{PrivKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{adec}}, \text{aenc}(x, k, y)).\text{out}(c_{k, \alpha(k)}^{\text{adec}}, x) \\ \mid_{k \in K^{\text{PrivKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{sign}}, x).\text{new } n.\text{out}(c_{k, \alpha(k)}^{\text{sign}}, \text{sign}(x, k, n)) \\ \mid_{k \in K^{\text{PubKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{check}}, \text{sign}(x, k, y)).\text{out}(c_{k, \alpha(k)}^{\text{check}}, x) \end{array} \\ \bar{Q} &= Q \mid \begin{array}{l} \mid_{k \in K^{\text{SymKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{senc}}, x).\text{new } n.\text{out}(c_{k, \alpha(k)}^{\text{senc}}, \text{senc}(x, \alpha(k), n)) \\ \mid_{k \in K^{\text{SymKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{sdec}}, \text{senc}(x, \alpha(k), y)).\text{out}(c_{k, \alpha(k)}^{\text{sdec}}, x) \\ \mid_{k \in K^{\text{PubKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{aenc}}, x).\text{new } n.\text{out}(c_{k, \alpha(k)}^{\text{aenc}}, \text{aenc}(x, \alpha(k), n)) \\ \mid_{k \in K^{\text{PrivKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{adec}}, \text{aenc}(x, \alpha(k), y)).\text{out}(c_{k, \alpha(k)}^{\text{adec}}, x) \\ \mid_{k \in K^{\text{PrivKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{sign}}, x).\text{new } n.\text{out}(c_{k, \alpha(k)}^{\text{sign}}, \text{sign}(x, \alpha(k), n)) \\ \mid_{k \in K^{\text{PubKey}}} !\text{in}(c_{k, \alpha(k)}^{\text{check}}, \text{check}(x, \alpha(k), y)).\text{out}(c_{k, \alpha(k)}^{\text{check}}, x), \end{array}\end{aligned}$$

where K^s denotes the keys of sort s of K . We call $\mathcal{T}_{\text{oracle}}$ the transformation taking a pair of protocols (P, Q) satisfying the aforementioned condition and returning the pair (\bar{P}, \bar{Q}) presently defined.

PROOF (sketch). First, thanks to Lemma 4.12, we know that P , \bar{P} , Q , and \bar{Q} disclose all their deducible keys.

(\Rightarrow) Given a witness of nonequivalence for $\bar{P} \approx_{\text{fwd}} \bar{Q}$, it is quite easy to build a witness of nonequivalence for $P \approx Q$ replacing the use of the oracle by the corresponding attacker construction. This yields a witness of nonequivalence for $P \approx Q$.

(\Leftarrow) This direction is actually more involved. The idea is to replace the use of an attacker construction, for example, an encryption with a deducible key, by the corresponding oracle. However, the attacker has the ability to use the same random seed more than once, whereas this is impossible when using the oracles to perform those computations. Thus, we first show that this additional ability does not give any power to the attacker. Then, we do the replacement as expected in order to conclude.

The full proof is provided in Appendix B.2. \square

Example 4.15. Continuing our example, this last transformation will add 10 branches (two per deducible key). For instance, regarding the key k_{AB} , the two following branches will be added:

For process P :

$$\begin{aligned} & !\text{in}(c_{k_{AB},k}^{\text{send}}, x). \text{new } n. \text{out}(c_{k_{AB},k}, \text{send}(x, k_{AB}, n)) \\ & | !\text{in}(c_{k_{AB},k}^{\text{sdec}}, \text{send}(x, k_{AB}, y)). \text{out}(c_{k_{AB},k}^{\text{sdec}}, x). \end{aligned}$$

For process Q :

$$\begin{aligned} & !\text{in}(c_{k_{AB},k}^{\text{send}}, x). \text{new } n. \text{out}(c_{k_{AB},k}, \text{send}(x, k, n)) \\ & | !\text{in}(c_{k_{AB},k}^{\text{sdec}}, \text{send}(x, k, y)). \text{out}(c_{k_{AB},k}^{\text{sdec}}, x). \end{aligned}$$

Regarding the keys pk_A, pk_B, pk_C , and sk_A , since $\alpha(k') = k'$ for each of these keys, we add the following branches on both sides:

$$\begin{aligned} & | !\text{in}(c_{pk_A, pk_A}^{\text{aenc}}, x). \text{new } n. \text{out}(c_{pk_A, pk_A}^{\text{aenc}}, \text{aenc}(x, pk_A, n)) \\ & | !\text{in}(c_{pk_B, pk_B}^{\text{aenc}}, x). \text{new } n. \text{out}(c_{pk_B, pk_B}^{\text{aenc}}, \text{aenc}(x, pk_B, n)) \\ & | !\text{in}(c_{pk_C, pk_C}^{\text{aenc}}, x). \text{new } n. \text{out}(c_{pk_C, pk_C}^{\text{aenc}}, \text{aenc}(x, pk_C, n)) \\ & | !\text{in}(c_{sk_C, sk_C}^{\text{adec}}, \text{aenc}(x, pk_C, y)). \text{out}(c_{sk_C, sk_C}^{\text{adec}}, x) \\ & | !\text{in}(c_{sk_C, sk_C}^{\text{sign}}, x). \text{new } n. \text{out}(c_{sk_C, sk_C}^{\text{sign}}, \text{sign}(x, sk_C, n)) \\ & | !\text{in}(c_{pk_A, pk_A}^{\text{check}}, \text{sign}(x, sk_A, y)). \text{out}(c_{pk_A, pk_A}^{\text{check}}, x) \\ & | !\text{in}(c_{pk_B, pk_B}^{\text{check}}, \text{sign}(x, sk_B, y)). \text{out}(c_{pk_B, pk_B}^{\text{check}}, x) \\ & | !\text{in}(c_{pk_C, pk_C}^{\text{check}}, \text{sign}(x, sk_C, y)). \text{out}(c_{pk_C, pk_C}^{\text{check}}, x). \end{aligned}$$

4.2.3. Transformation T_{fwd} . Thanks to Lemmas 4.12 and 4.14, we are now able to formally define our transformation that gets rid of a fully active attacker. For every pair of protocols (P, Q) in \mathcal{C}_{pp} , we consider

$$T_{\text{fwd}}(P, Q) = \{T_{\text{oracle}}(P', Q') \mid (P', Q') \in T_{\text{key}}(P, Q)\}.$$

Combination of the two previous results yields the desired result.

PROPOSITION 4.16. *Let P and Q be two protocols in \mathcal{C}_{pp} . We have that*

$$P \approx Q \text{ if and only if } P' \approx_{\text{fwd}} Q' \text{ for some } (P', Q') \in T_{\text{fwd}}(P, Q).$$

5. ENCODING PROTOCOLS INTO REAL-TIME GPDAS

We first introduce the notion of a real-time generalized pushdown automaton (GPDA) (see Section 5.1) before explaining in detail (see Sections 5.2 and 5.3) our encoding from protocols to real-time generalized pushdown automata. More precisely, for any process $P \in \mathcal{C}_{\text{pp}}$, we show that it is possible to define a polynomial-sized real-time generalized pushdown automaton \mathcal{A}_P such that trace equivalence w.r.t. the forwarder semantics coincides with language equivalence of the two corresponding automata.

We illustrate the different steps of our translation of protocols to automata using a (mock) ping-pong protocol P_{toy} . We define $\text{io}(c, R, w) \stackrel{\text{def}}{=} \text{in}(c, R).\text{out}(c, w)$.

$$P_{\text{toy}} = \begin{array}{l} | \text{in}(c_1, \text{start}).\text{new } r_1.\text{out}(c_1, \text{senc}(a, k_2, r_1)) \\ | \text{! in}(c_2, \text{senc}(x, k_2, z_1)).\text{new } r_1.\text{out}(c_2, \text{senc}(x, k_1, r_1)) \\ | \text{! in}(c_3, x).\text{new } r_1.\text{out}(c_3, \text{senc}(x, k_2, r_1)) \\ | \text{! in}(c_4, \text{senc}(\text{senc}(x, k_1, z_1), k_2, z_2)).\text{new } r_1, r_2.\text{out}(c_4, \text{senc}(\text{senc}(x, k_2, r_1), k_1, r_2)) \\ | \text{! in}(c_5, \text{senc}(\text{senc}(x, k_2, z_1), k_1, z_2)).\text{out}(c_5, x) \end{array}$$

For illustrative purpose, we consider different execution traces of this protocol. For instance, we have that $(\text{tr}_1, \sigma_1) \in \text{trace}_{\text{fwd}}(P_{\text{toy}})$ where:

$$\begin{array}{l} - \text{tr}_1 = \text{io}(c_1, \text{start}, w_1).\text{io}(c_2, w_1, w_2).\text{io}(c_3, w_2, w_3).\text{io}(c_4, w_3, w_4).\text{io}(c_5, w_4, w_5), \text{ and} \\ - \sigma_1 = \{w_1 \triangleright \text{senc}(a, k_2, r_1), w_2 \triangleright \text{senc}(a, k_1, r_2), w_3 \triangleright \text{senc}(\text{senc}(a, k_1, r_2), k_2, r_3), \\ \quad w_4 \triangleright \text{senc}(\text{senc}(a, k_2, r_4), k_1, r_5), w_5 \triangleright a\}. \end{array}$$

This execution may be continued as follows:

$$\begin{array}{l} - \text{tr}_2 = \text{io}(c_3, w_1, w_6).\text{io}(c_2, w_6, w_7).\text{io}(c_5, w_7, w_8), \text{ and} \\ - \sigma_2 = \{w_6 \triangleright \text{senc}(\text{senc}(a, k_2, r_1), k_2, r_6), w_7 \triangleright \text{senc}(\text{senc}(a, k_2, r_1), k_1, r_7), w_8 \triangleright a\}. \end{array}$$

Let $\sigma_{1/2} = \sigma_1 \cup \sigma_2$. We have that $(\text{tr}_1.\text{tr}_2, \sigma_{1/2})$ is a trace of P_{toy} w.r.t. the forwarder semantics. We have that the test $w_5 = w_8$ is valid in $\sigma_{1/2}$. Indeed $w_5\sigma_{1/2}\downarrow = w_8\sigma_{1/2}\downarrow = a$.

We have also that $(\text{tr}'_1, \sigma'_1) \in \text{trace}_{\text{fwd}}(P_{\text{toy}})$ with:

$$\begin{array}{l} - \text{tr}'_1 = \text{io}(c_1, \text{start}, w_1).\text{io}(c_2, w_1, w_2).\text{io}(c_3, w_2, w_3).\text{io}(c_4, w_3, w_4).\text{io}(c_3, w_4, w_5), \text{ and} \\ - \sigma'_1 = \{w_1 \triangleright \text{senc}(a, k_2, r_1), w_2 \triangleright \text{senc}(a, k_1, r_2), w_3 \triangleright \text{senc}(\text{senc}(a, k_1, r_2), k_2, r_3), \\ \quad w_4 \triangleright \text{senc}(\text{senc}(a, k_2, r_4), k_1, r_5), w_5 \triangleright \text{senc}(\text{senc}(\text{senc}(a, k_2, r_4), k_1, r_5), k_2, r_6)\}. \end{array}$$

This execution may be continued as follows:

$$\begin{array}{l} - \text{tr}'_2 = \text{io}(c_4, w_5, w_6).\text{io}(c_5, w_6, w_7).\text{io}(c_4, w_5, w_8).\text{io}(c_5, w_8, w_9), \\ - \sigma'_2 = \{w_6 \triangleright \text{senc}(\text{senc}(\text{senc}(a, k_2, r_4), k_2, r_7), k_1, r_8), w_7 \triangleright \text{senc}(a, k_2, r_4), \\ \quad w_8 \triangleright \text{senc}(\text{senc}(\text{senc}(a, k_2, r_4), k_2, r_7), k_1, r_8), w_9 \triangleright \text{senc}(a, k_2, r_4)\}. \end{array}$$

Let $\sigma'_{1/2} = \sigma'_1 \cup \sigma'_2$. We have that $(\text{tr}'_1.\text{tr}'_2, \sigma'_{1/2})$ is a trace of P_{toy} w.r.t. the forwarder semantics. We have that the test $w_7 = w_9$ is valid in $\sigma'_{1/2}$.

Fig. 3. Running example.

THEOREM 5.1. *Let P and Q in \mathcal{C}_{pp} . We have that*

$$P \approx_{\text{fwd}} Q \iff \mathcal{L}(\mathcal{A}_P) = \mathcal{L}(\mathcal{A}_Q).$$

The proof of this theorem consists of three main steps:

- (1) First, we provide a new characterization of trace equivalence w.r.t. the forwarder semantics. Intuitively, we show that it is not necessary to consider all possible tests (when checking static equivalence). Indeed, our Lemma 5.8 states that it is sufficient to check for constant tests (i.e., tests of the form $x = c$, where c is a constant) and some specific class of tests that we call *guarded* and *pulled up*.
- (2) Then we associate to processes $P, Q \in \mathcal{C}_{\text{pp}}$ real-time GPDAs that check whether they satisfy the same constant tests (Lemma 5.9).
- (3) Then we associate to processes $P, Q \in \mathcal{C}_{\text{pp}}$ real-time GPDAs that check whether they satisfy the same guarded tests (Lemma 5.11).

Throughout this section, we illustrate the definitions with the protocol displayed in Figure 3. This example should be read step by step when reading the examples of this section.

5.1. Generalized Pushdown Automata

Language equivalence of deterministic pushdown automata (DPA) is known to be decidable [Sénizergues 2001]. We actually encode equivalence of protocols into a fragment of DPA: real-time GPDA with final-state acceptance. GPDAs differ from deterministic pushdown automata (DPA) as they can unstack several symbols at a time. Real-time automata are automata that do not include epsilon-transitions. Formally, the class of real-time GPDA is defined as follows:

Definition 5.2. A real-time GPDA is a 7-tuple $\mathcal{A} = (Q, \Pi, \Gamma, q_0, \omega, Q_f, \delta)$, where Q is a finite set of states, $q_0 \in Q$ is an initial state, $Q_f \subseteq Q$ is a set of accepting states, Π is a finite input alphabet, Γ is a finite stack alphabet, ω is the initial stack symbol, and $\delta : (Q \times \Pi \times \Gamma_0) \rightarrow Q \times \Gamma_0$ is a partial transition function such that

- Γ_0 is a finite subset of Γ^* ; and
- for any $(q, a, x) \in \text{dom}(\delta)$ and y suffix strict of x , we have that $(q, a, y) \notin \text{dom}(\delta)$.

Let $q, q' \in Q$, $u, u', \gamma \in \Gamma^*$, $m \in \Pi^*$, $a \in \Pi$; we note $(qu\gamma, am) \rightsquigarrow_{\mathcal{A}} (q'u' \epsilon, m)$ if $(q', u') = \delta(q, a, \gamma)$. The relation $\rightsquigarrow_{\mathcal{A}}^*$ is the reflexive and transitive closure of $\rightsquigarrow_{\mathcal{A}}$. For every $qu, q'u'$ in $Q\Gamma^*$ and $m \in \Pi^*$, we note $qu \xrightarrow{m}_{\mathcal{A}} q'u'$ if and only if $(qu, m) \rightsquigarrow_{\mathcal{A}}^* (q'u' \epsilon, \epsilon)$. For the sake of clarity, a transition from q to q' reading a , popping γ from the stack, and pushing u' will be denoted by $q \xrightarrow{a; \gamma/u'} q'$.

Let \mathcal{A} be a GPDA. The language recognized by \mathcal{A} is defined by

$$\mathcal{L}(\mathcal{A}) = \{m \in \Pi^* \mid q_0 \omega \text{start} \xrightarrow{m}_{\mathcal{A}} q_f u \text{ for some } q_f \in Q_f \text{ and } u \in \Gamma^*\}.$$

Note that the language is defined starting with the word ωstart in the stack.

A real-time GPDA can easily be converted into a DPA by adding new states and ϵ -transitions. Thus, the problem of language equivalence for two real-time GPDAs \mathcal{A}_1 and \mathcal{A}_2 , that is, deciding whether $\mathcal{L}(\mathcal{A}_1) = \mathcal{L}(\mathcal{A}_2)$, is decidable [Sénizergues 2001]. Whether deciding equivalence of real-time GPDA could be easier than deciding equivalence of DPA is an open question.

5.2. Characterization of Trace Equivalence

To construct the automaton associated to a process $P \in \mathcal{C}_{pp}$, we need to construct an automaton that recognizes any execution of P and the corresponding valid tests.

We first propose a new characterization of trace equivalence allowing us to restrict our attention to executions of P and valid tests that have a special form.

Given an execution trace (tr, σ) and an element w of a frame σ , we can extract from tr the sequence of actions that were conducted for the production of this element w .

Definition 5.3. Let P be a protocol in \mathcal{C}_{pp} ; tr be a trace of P w.r.t. the forwarder semantics, that is, such that $(\text{tr}, \sigma) \in \text{trace}_{\text{fwd}}(P)$ for some σ ; and w be a variable that occurs in tr . The *sequence associated to w in tr* , denoted $\text{seq}_{\text{tr}}(w)$, is the subsequence of tr of the following form:

$$\text{seq}_{\text{tr}}(w) = \text{io}(c_{i_0}, \text{start}, w_{j_0}).\text{io}(c_{i_1}, w_{j_0}, w_{j_1}) \dots \text{io}(c_{i_p}, w_{j_{p-1}}, w).$$

Example 5.4. Consider the protocol defined in Figure 3. Then,

- $\text{seq}_{\text{tr}_1, \text{tr}_2}(w_5) = \text{tr}_1$;
- $\text{seq}_{\text{tr}_1, \text{tr}_2}(w_8) = \text{io}(c_1, \text{start}, w_1).\text{tr}_2$;
- $\text{seq}_{\text{tr}_1, \text{tr}_2}(w_7) = \text{tr}'_1.\text{io}(c_4, w_5, w_6).\text{io}(c_5, w_6, w_7)$; and
- $\text{seq}_{\text{tr}'_1, \text{tr}'_2}(w_9) = \text{tr}'_1.\text{io}(c_4, w_5, w_8).\text{io}(c_5, w_8, w_9)$.

We consider some particular class of tests, called *pulled-up* tests.

Definition 5.5. Let P be a protocol in \mathcal{C}_{pp} , $(tr, \sigma) \in \text{trace}_{\text{fwd}}(P)$, and $w, w' \in \text{dom}(\sigma)$ such that

- (1) the test $w = w'$ is σ -valid, that is, $w\sigma = w'\sigma$; and
- (2) the test $w = w'$ is σ -guarded; that is, the head symbol of $w\sigma$ (or equivalently $w'\sigma$) is in $\{\text{senc}, \text{aenc}, \text{sign}\}$.

Let $\text{io}(c_{i_0}, \text{start}, w_{j_0}) \dots \text{io}(c_{i_p}, w_{j_{p-1}}, w_{j_p})$ be the maximal common prefix of $\text{seq}_{tr}(w)$ and $\text{seq}_{tr}(w')$. The test $w = w'$ is said to be *pulled up* in (tr, σ) if $p = 0$, or $p \geq 1$ and $w\sigma$ does not occur as a subterm in $w_{j_0}\sigma, \dots, w_{j_{p-1}}\sigma$.

Intuitively, to perform a test $w = w'$, the attacker (who acts as a forwarder) relies on the protocol rules to produce successive outputs, and ultimately the ones stored in w and w' . The attacker may produce w and w' independently (the common prefix of $\text{seq}_{tr}(w)$ and $\text{seq}_{tr}(w')$ is empty), and in such a case the test is pulled up by definition. This is not, of course, always possible. In particular, a test $w = w'$ satisfying conditions (1) and (2) of the previous definition is necessarily a “forked” test, that is, a test for which the common prefix of $\text{seq}_{tr}(w)$ and $\text{seq}_{tr}(w')$ is not reduced to the empty sequence, and thus $p \geq 1$. Indeed, $w\sigma$ is a term of the form $f(u, k, r)$ with some random r . Since nonces are uniquely generated, the variable w_i that generates it, that is, the smallest i such that r occurs in $w_i\sigma$, occurs both in $\text{seq}_{tr}(w)$ and $\text{seq}_{tr}(w')$. For this kind of “forked” test, we can restrict the attacker to consider tests that are pulled up; that is, we consider tests for which the size of the common prefix between $\text{seq}_{tr}(w)$ and $\text{seq}_{tr}(w')$ is reduced to the minimum. This can be done by duplicating some execution steps since all the branches are under a replication.

Example 5.6. Continuing our running example, we have that $w_5 = w_8$ is a test that is $\sigma_{1/2}$ -valid but it is not $\sigma_{1/2}$ -guarded since $w_5\sigma_{1/2} = w_8\sigma_{1/2} = a$.

The test $w_7 = w_9$ is a test that is $\sigma'_{1/2}$ -valid and $\sigma'_{1/2}$ -guarded. Indeed, we have that $w_7\sigma'_{1/2} = w_9\sigma'_{1/2} = \text{senc}(a, k_2, r_4)$. The maximal common prefix of $\text{seq}_{tr_1, tr_2}(w_7)$ and $\text{seq}_{tr_1, tr_2}(w_9)$ is actually

$$tr'_1 = \text{io}(c_1, \text{start}, w_1). \text{io}(c_2, w_1, w_2). \text{io}(c_3, w_2, w_3). \text{io}(c_4, w_3, w_4). \text{io}(c_3, w_4, w_5).$$

Actually, $w_7\sigma'_{1/2}$ occurs as a subterm in $w_4\sigma'_{1/2}$, and thus the test $w_7 = w_9$ is not pulled up in $(tr'_1, tr'_2, \sigma'_{1/2})$.

We are now able to state our characterization lemma. Intuitively, we show that for tests that are valid and guarded, it is sufficient to consider pulled-up tests. We first illustrate through an example how a test that is valid and guarded can be converted into a pulled-up one.

Example 5.7. Continuing Example 5.4, we consider the test $w_7 = w_9$, which is not pulled up in $(tr'_1, tr'_2, \sigma'_{1/2})$. Consider the execution

$$tr' = tr'_1. \text{io}(c_4, w_5, w_6). \text{io}(c_5, w_6, w_7). \text{io}(c_3, w_4, w_8). \text{io}(c_4, w_8, w_9). \text{io}(c_5, w_9, w_{10}).$$

This execution is almost similar to tr'_1, tr'_2 . The main difference is that the computation performed at the end of tr'_1 using channel c_3 with input w_4 is duplicated. Both $\text{io}(c_3, w_4, w_5)$ and $\text{io}(c_3, w_4, w_8)$ occur in tr' . The resulting frame is

$$\begin{aligned} \sigma'_1 \cup \{ & w_6 \triangleright \text{senc}(\text{senc}(\text{senc}(a, k_2, r_4), k_2, r_7), k_1, r_8), \quad w_7 \triangleright \text{senc}(a, k_2, r_4), \\ & w_8 \triangleright \text{senc}(\text{senc}(\text{senc}(a, k_2, r_4), k_1, r_5), k_2, r'_6), \\ & w_9 \triangleright \text{senc}(\text{senc}(\text{senc}(a, k_2, r_4), k_2, r'_7), k_1, r'_8), \quad w_{10} \triangleright \text{senc}(a, k_2, r_4) \}. \end{aligned}$$

The terms stored in w_5 and w_8 differ by their random seeds:

$\text{senc}(\text{senc}(\text{senc}(a, k_2, r_4), k_1, r_5), k_2, r_6)$ and $\text{senc}(\text{senc}(\text{senc}(a, k_2, r_4), k_1, r_5), k_2, r'_6)$.

This frame is almost the same as $\sigma'_{1/2}$ with an additional element (w_8). The term stored in w_8 is the same as the one stored in w_5 up to the choice of some random seeds (r_6 is replaced by the fresh random r'_6). Moreover, the presence of this additional element leads us to reindex the following elements of the frame and to replace some occurrences of r_6 with r'_6 . It is important to note that the introduced randoms r'_6 and r'_8 could potentially break equality tests. They, however, do not appear anymore in the last outputted term stored in w_{10} that is checked for equality.

This example shows that when considering the trace $(\text{tr}'_1, \text{tr}'_2, \sigma'_{1/2})$, we may have to consider the test $w_7 = w_9$, which is not pulled up. However, this test is essentially the same as the pulled-up test $w_7 = w_{10}$ issued from the trace given earlier.

The transformation explained in the previous example can be generalized to any protocol.

LEMMA 5.8. *Let P and Q be two protocols in \mathcal{C}_{pp} , and then $P \approx_{\text{fwd}} Q$ if and only if the following four conditions are satisfied:*

- CONST_P**: *For all $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$, there exists a frame σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$ and for every $w, w' \in \text{dom}(\sigma_P)$ and for every constant $c \in \Sigma_0 \cup \{\text{start}\}$, $w\sigma_P = w'\sigma_Q = c$ if and only if there exists a constant $c' \in \Sigma_0 \cup \{\text{start}\}$ such that $w\sigma_Q = w'\sigma_Q = c'$.*
- CONST_Q**: *Similarly swapping the roles of P and Q .*
- GUARDED_P**: *For all $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$, there exists a frame σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$ and every test that is σ_P -valid, σ_P -guarded, and pulled up in (tr, σ_P) is also σ_Q -valid, σ_Q -guarded, and pulled up in (tr, σ_Q) .*
- GUARDED_Q**: *Similarly swapping the roles of P and Q .*

PROOF (sketch). (\Rightarrow) For this direction, when considering **CONST_P**, the only difficulty is to show that the test $w\sigma_Q = w'\sigma_Q$ leads to a constant c' . Actually, such a test cannot lead to a guarded test since otherwise a replay of the entire sequence (this replay is possible since we consider a class of protocol that allows this) will lead to a different guarded term in Q and not in P (due to the presence of fresh randoms in guarded terms).

When considering **GUARDED_P**, the difficulty is to show that the test $w = w'$ is necessarily pulled up in (tr, σ_Q) . Let $\text{pref} = \text{io}(c_{i_0}, \text{start}, w_{j_0}) \dots \text{io}(c_{i_p}, w_{j_{p-1}}, w_{j_p})$ be the maximal common prefix of $\text{seq}_{\text{tr}}(w)$ and $\text{seq}_{\text{tr}}(w')$. Since $w = w'$ is pulled up in (tr, σ_P) , we know that the first occurrence of $w\sigma_P$ in $\text{pref}\sigma_P$ is at the very end of the sequence. We can easily show that $w = w'$ is σ_Q -valid and σ_Q -guarded, and thus $w\sigma_Q$ occurs also as a subterm in $\text{pref}\sigma_Q$. The only problem is if $w\sigma_Q$ occurs in $\text{pref}\sigma_Q$ but not at the very end of this sequence. The idea is that in such a case, we can modify the trace (tr, σ_Q) and the test $w = w'$ to build $(\text{tr}^*, \sigma_Q^*)$ and a new test $w_* = w'_*$, which will be pulled up in $(\text{tr}^*, \sigma_Q^*)$. The idea is to split the two sequences $\text{seq}_{\text{tr}}(w)$ and $\text{seq}_{\text{tr}}(w')$ earlier without compromising the fact that the test will be valid in the resulting frame. This corresponds to the construction illustrated in Example 5.7. This trace tr^* is actually a witness of nonequivalence. Actually, the test $w_* = w'_*$ is *a fortiori* not valid on the P side, and this contradicts our hypothesis $P \approx_{\text{fwd}} Q$.

(\Leftarrow) Actually, for this direction, assume that we have a witness of the fact that $P \not\approx Q$, that is, a trace $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$, a trace $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$, and a test $w = w'$ that is σ_P -valid but not σ_Q -valid. In case the resulting term is a constant, we easily conclude that **CONST_P** fails. Otherwise, it means that $w = w'$ is σ_P -guarded. In order to show that

GUARDED_P fails, we have to ensure that the test $w = w'$ is pulled up w.r.t. (tr, σ_P) . Since this is not necessarily the case, we have to build another trace $(\text{tr}^*, \sigma_P^*)$ that will lead us to a pulled-up test. Roughly, the transformation consists of splitting the two sequences $\text{seq}_{\text{tr}}(w)$ and $\text{seq}_{\text{tr}}(w')$ earlier without compromising the fact that the test will be valid in the resulting frame. Actually, such a transformation cannot transform a test that was not valid in a valid one, and thus this test is still not valid for Q and it is still a witness of nonequivalence, but a pulled-up one allowing us to conclude. \square

The detailed proof can be found in Appendix C.1.

5.3. From Trace Equivalence to Language Equivalence

Our goal is to associate an automaton \mathcal{A}_P to a protocol P such that \mathcal{A}_P recognizes the words (a sequence of channels) that correspond to a possible execution of the protocol. The stack of the automaton \mathcal{A}_P is used to store a (partial) representation of the last outputted term. This first requires us to convert a term into a word.

Given an input term or an output term u (see Section 3.1), we define inductively \bar{u} in the following way:

$$\begin{cases} \bar{u} = \bar{v}.k & \text{if } u = f(v, k, r) \text{ and } f \in \{\text{senc}, \text{aenc}, \text{sign}\} \\ \bar{c} = \omega c & \text{for any constant } c \in \Sigma_0 \cup \{\text{start}\} \\ \bar{x} = \epsilon & \text{for any variable } x, \end{cases}$$

where ϵ denotes the empty word. Note that, using this representation, random seeds are not part of the encoding. We denote by $\|u\|$ the *height* of the term u , which is equal to the number of occurrences of senc , aenc , and sign in u .

We now consider an arbitrary ping-pong protocol P (using the same notation as the one introduced in Section 3):

$$P \stackrel{\text{def}}{=} \bigg|_{i=1}^n \bigg|_{j=1}^{p_i} \text{lin}(c_i, u_i^j). \text{new } r_1. \dots \text{new } r_{k_i^j}. \text{out}(c_i, v_i^j). \quad (*)$$

In the remainder of the section, we denote by Σ_0^P the finite set of constants of $\Sigma_0 \cup \{\text{start}\}$ that actually occur in the protocol P .

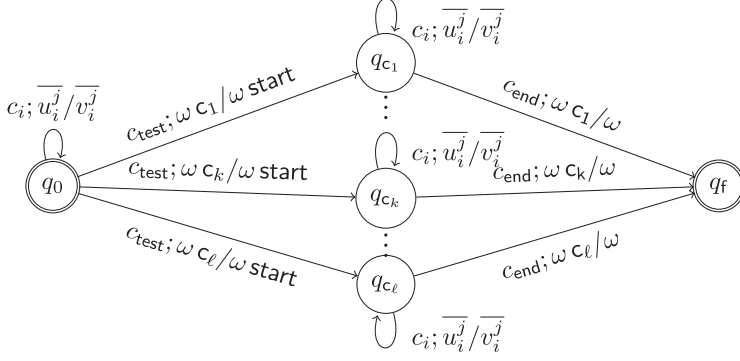

5.3.1. Encoding of the Conditions CONST_P and CONST_Q . We first build an automaton that recognizes tests of the form $w = w'$ such that the corresponding term is actually a constant. We define $\mathcal{A}_{\text{CONST}}^P$ as follows:

$$\mathcal{A}_{\text{CONST}}^P = (\{q_0, q_t\} \cup \{q_c \mid c \in \Sigma_0^P\}, \{c_1, \dots, c_n\} \cup \{c_{\text{test}}, c_{\text{end}}\}, \Sigma_0^P, q_0, \omega, \{q_0, q_t\}, \delta),$$

where the transition function δ is defined as follows:

- (1) for every $q \in \{q_0\} \cup \{q_c \mid c \in \Sigma_0^P\}$, for every $i \in \{1, \dots, n\}$, and for every $j \in \{1, \dots, p_i\}$,
there is a transition $q \xrightarrow{c_i; \bar{u}_j^i / \bar{v}_j^i} q;$
- (2) for every constant c , there is a transition $q_0 \xrightarrow{c_{\text{test}}; \omega c / \omega \text{start}} q_c;$ and
- (3) for every constant c , there is a transition $q_c \xrightarrow{c_{\text{end}}; \omega c / \omega} q_t.$

The automaton is depicted in Figure 4. Intuitively, the basic building blocks (e.g., q_0 with the transitions from q_0 to itself) mimic an execution of P where each input is fed with the last outputted term. Then, to recognize the tests of the form $w = w'$ that are true in such an execution, it is sufficient to memorize the constant c that is associated to w (adding a new state q_c) and to see whether it is possible to reach a state where the stack contains c again. More formally, we have the following result.

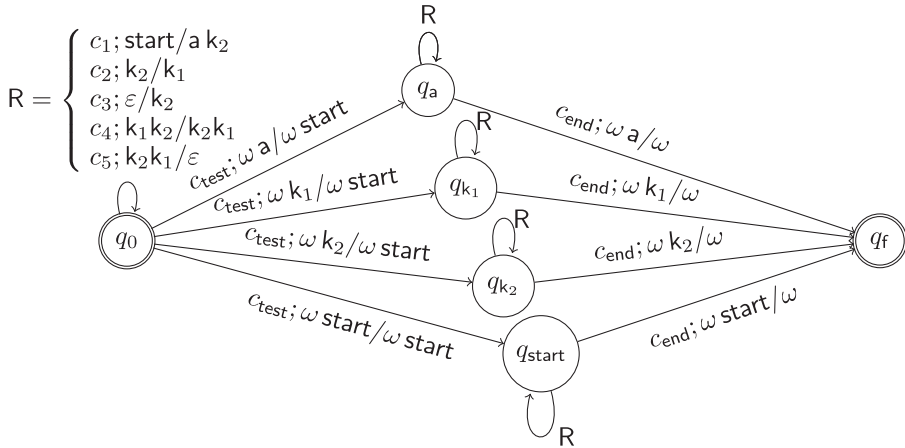
Fig. 4. Automaton $\mathcal{A}_{\text{CONST}}^P$.Fig. 5. Automaton $\mathcal{A}_{\text{GUARDED}}^P$.

LEMMA 5.9. Let P and Q be two protocols in \mathcal{C}_{pp} ; the two real-time GPDA's $\mathcal{A}_{\text{CONST}}^P$ and $\mathcal{A}_{\text{CONST}}^Q$ are such that

P and Q satisfy conditions CONST_P and CONST_Q if and only if
 $\mathcal{L}(\mathcal{A}_{\text{CONST}}^P) = \mathcal{L}(\mathcal{A}_{\text{CONST}}^Q)$.

The proof can be found in Appendix C.2.

Example 5.10. Going back to our running example, that is, the protocol P described in Figure 3, the automaton $\mathcal{A}_{\text{CONST}}^P$ is depicted as follows:



The word that represents the trace $(\text{tr}_1.\text{tr}_2, \sigma_{1/2})$ and the test $w_5 = w_8$ as given in Figure 3 is $c_1c_2c_3c_4c_5c_{\text{test}}c_1c_3c_2c_5c_{\text{end}}$. The fact that this test is a valid one that leads to a constant a means that the word will be accepted by the automaton given previously. The corresponding run goes through the state q_a and halts in state q_f .

$\mathcal{A}_{\text{CONST}}^P$ has a number of states polynomial in the number of constants in P and for each state a number of transitions linear in the number of branches in P . Thus, $\mathcal{A}_{\text{CONST}}^P$ is of size polynomial with respect to the size of P .

5.3.2. Encoding of the Conditions GUARDED_P and GUARDED_Q. Capturing tests that lead to nonconstant symbols (i.e., terms of the form $f(u, k, r)$ with $f \in \{\text{senc}, \text{aenc}, \text{sign}\}$) is more tricky for several reasons. First, it is not possible anymore to memorize the resulting term in a state of the automaton. Second, names of sort rand play a role in such a test, while they are forgotten in our encoding. We rely on our characterization introduced in Section 5.2, and we construct a more complex automaton that uses some special track symbols to encode when randomized ciphertexts may be reused.

More precisely, we consider

- $\Pi = \{c_1, \dots, c_n, c_{\text{test}}, c_{\text{end}}\} \cup \{c_{\text{fork}}^i \mid 1 \leq i \leq n\}$, and
- $\Gamma = \Sigma_0^P \cup \{\text{test}\} \cup \{(\text{fork}_i^j, k) \mid 1 \leq i \leq n, 1 \leq j \leq p_i, \text{ and } 1 \leq k \leq \|u_i^j\|\}$.

Note that n and p_i are induced by the definition of protocol P (see Equation (*)). The input alphabet contains the channel names c_1, \dots, c_n , plus some additional symbols, denoted $c_{\text{fork}}^1, \dots, c_{\text{fork}}^n$, that will be used once and whose purpose will be to mark the end of the common prefix between $\text{seq}_{\text{tr}}(w)$ and $\text{seq}_{\text{tr}}(w')$.

The stack alphabet is more involved. We still have one symbol per constant in Σ_0^P and a special symbol test that will be put on top of the stack when the stack contains the target term (i.e., $w\sigma$). In such an automaton, the idea is to consider pulled-up tests only. The tile (fork_i^j, k) is placed on the stack when the automaton has finished to build the term corresponding to the left-hand side of a pulled-up test.

The transition function δ is defined as follows:

- (1) for $q \in \{q_0, q_1, q_2\}$, for every $i \in \{1, \dots, n\}$ and $j \in \{1, \dots, p_i\}$, there is a transition $q \xrightarrow{c_i; \overline{u_i^j}/\overline{v_i^j}} q$;
- (2) for every $i \in \{1, \dots, n\}$ and $j \in \{1, \dots, p_i\}$ such that $\|v_i^j\| \geq 1$, there is a transition

$$q_0 \xrightarrow{c_{\text{fork}}^i; \overline{u_i^j}/\omega \overline{v_i^j} (\text{fork}_i^j, \|v_i^j\|)} q_1;$$

- (3) for every $i \in \{1, \dots, n\}$ and $j \in \{1, \dots, p_i\}$, for every $i' \in \{1, \dots, n\}$ and $j' \in \{1, \dots, p_{i'}\}$, for every m such that $1 < m \leq \|v_{i'}^{j'}\|$, and for every subterm u_0 of $u_{i'}^{j'}$ of height $k \in \{1, \dots, m-1\}$ such that $\overline{u_i^j} = \overline{u_0}.s'$, there is a transition

$$q_1 \xrightarrow{c_i; \overline{u_0}.(\text{fork}_{i'}^{j'}, m).s' / (\text{fork}_{i'}^{j'}, m-k)\overline{v_{i'}^{j'}}} q_1;$$

- (4) for every $i \in \{1, \dots, n\}$ and $j \in \{1, \dots, p_i\}$, for every m such that $1 \leq m \leq \|v_i^j\|$, there is a transition $q_1 \xrightarrow{c_{\text{test}}; (\text{fork}_i^j, m)/\text{test } s_{i,m}^j} q_2$, where $s_{i,m}^j$ is the suffix of length $\|v_i^j\| - m$ of $\overline{v_i^j}$; and
- (5) there is a transition $q_2 \xrightarrow{c_{\text{end}}; \text{test}/\epsilon} q_f$.

The loop in q_0 (item 1) represents the regular execution of the protocol by the attacker: through unstacking and stacking, he or she builds a term on the stack along a particular

trace. The transitions $q_0 \xrightarrow{c_{\text{fork}}^i; z/z'} q_1$ (item 2) enable him or her to mark a fork when building a test in his or her frame with a particular stack symbol fork_i^j , enriched with some information. Intuitively, the part of the execution that is performed until here should correspond to the maximal prefix shared between the sequences $\text{seq}_{\text{tr}}(w)$ and $\text{seq}_{\text{tr}}(w')$. By looping in q_1 , the attacker can continue building the first term of an equality, following the usual execution of the protocol, if it were for the presence of the stack symbol (fork_i^j, k) , which can only go down on the stack for at most $k - 1$ times. When the symbol (fork_i^j, k) appears on top of the stack, the attacker may decide that he or she has built the first part of a pulled-up test. Then test will be put on the top of the stack and a part of the stack (following the instructions memorized in the symbol (fork_i^j, k)) will be regenerated. The idea is that the stack has to contain the same term as the one stored just after forking. Then the attacker tries to build the second member of the test. If this second term manages to end up exactly as the previous one (the position in the stack is marked using the tile test), an equality is reached and the word is recognized by the automata, witnessing the equality induced by the pulled-up test.

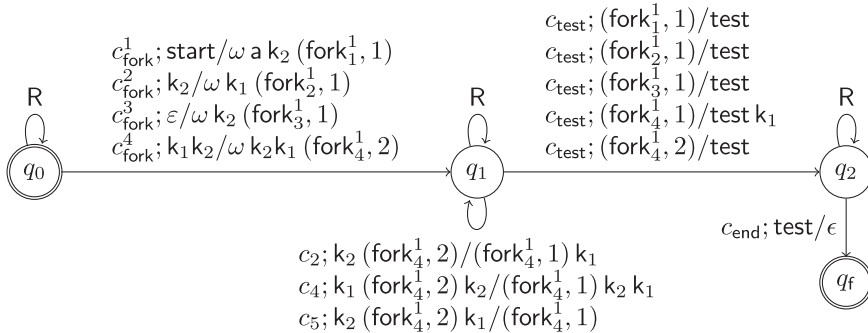
What remains now is to prove that P and Q satisfy conditions GUARDED_P and GUARDED_Q if and only if $\mathcal{L}(\mathcal{A}_{\text{GUARDED}}^P) = \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q)$. This is formally stated in the following lemma.

LEMMA 5.11. *Let P and Q be two protocols in \mathcal{C}_{pp} ; the two real-time GPDA $\mathcal{A}_{\text{GUARDED}}^P$ and $\mathcal{A}_{\text{GUARDED}}^Q$ are such that*


$$P \text{ and } Q \text{ satisfy conditions } \text{GUARDED}_P \text{ and } \text{GUARDED}_Q \text{ if and only if } \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^P) = \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q).$$

The proof can be found in Appendix C.2.

Example 5.12. Going back to our running example, that is, the protocol P described in Figure 3, the automaton $\mathcal{A}_{\text{GUARDED}}^P$ is depicted as follows:



The set of transitions R is the one defined in Example 5.10. The situation where the stack symbol (fork_i^j, k) goes down, for instance, occurs when considering the word $c_1 c_2 c_3 c_{\text{fork}}^4 c_3 c_4 c_5 c_{\text{test}} c_3 c_4 c_5$. The evolution of the stack during the run of the automaton is depicted as follows. On the second line, we can see that this symbol goes down and k goes from 2 to 1:



The trace $(\text{tr}, \sigma) \in \text{trace}_{\text{fwd}}(P)$ and the pulled-up test $w = w'$ that correspond to this execution are the ones introduced in Example 5.7, that is, tr' together with the test $w_7 = w_{10}$.

We can notice that up to the special stack symbols, namely, test and (fork_i^j, k) , the contents of the stack after reading c_{fork}^i (here $i = 4$) and c_{test} are the same. The stack actually represents the term obtained after executing the common prefix shared between $\text{seq}_{\text{tr}}(w_7)$ and $\text{seq}_{\text{tr}}(w_{10})$, that is, $\text{senc}(\text{senc}(a, k_2, r_4), k_1, r_5)$ stored in w_4 . We have also that the contents of the stack before reading c_{test} and after reading c_{end} are also the same (up to some special stack symbols). They actually represent the terms stored respectively in w_7 and w_{10} .

Note that $\mathcal{A}_{\text{GUARDED}}^P$ has a fixed number of states and a polynomial number of transitions: transitions are added for each branch and suffix of any input term. Thus, $\mathcal{A}_{\text{GUARDED}}^P$ is of size polynomial with respect to the size of P .

6. FROM LANGUAGE EQUIVALENCE TO TRACE EQUIVALENCE

We have seen how to encode trace equivalence between processes in \mathcal{C}_{pp} into language equivalence between real-time GPDA. The two problems are actually *equivalent*. Indeed, in this section, we show that we can conversely encode any real-time GPDA \mathcal{A} into a process $P_{\mathcal{A}}$ in \mathcal{C}_{pp} such that $\mathcal{L}(\mathcal{A}) \subseteq \mathcal{L}(\mathcal{B})$ implies $P_{\mathcal{A}} \sqsubseteq P_{\mathcal{B}}$.

Consider an automaton $\mathcal{A} = (Q, \Pi, \Gamma, q_0, \omega, Q_f, \delta)$. The process $P_{\mathcal{A}}$ associated to \mathcal{A} is built using symmetric encryption only. For the purpose of the encoding, we consider the following constants of sort SymKey :

- for each $q \in Q$, we denote q its counterpart in Σ_0 ;
- for each $\alpha \in \Gamma$, we denote k_{α} its counterpart in Σ_0 ; and
- a constant k_{well} .

Let also c_0, c_a, c_f with $a \in \Pi$ be constant symbols of sort channel in \mathcal{Ch} . Words in Γ^* , that is, stacks, will be represented through nested encryptions with keys representing their counterparts in Γ . For the sake of brevity, given a word $u = \alpha_1 \dots \alpha_p$ of Γ^* , we denote by $\overline{x.u}$

- either the term $\text{senc}(\dots \text{senc}(x, k_{\alpha_1}, z_1) \dots, k_{\alpha_p}, z_p)$, where z_1 through z_p are variables used for nonces when $\overline{x.u}$ is used in as an input pattern;
- or the term $\text{senc}(\dots \text{senc}(x, k_{\alpha_1}, r_1) \dots, k_{\alpha_p}, r_p)$, where r_1 through r_p are fresh randoms when $\overline{x.u}$ is used as an output pattern.

Next, we use $\text{new}\tilde{r}$ as a shortcut for $\text{new}r_1 \dots \text{new}r_p$ such that the sequence will bind every nonce occurring in the following output.

The stack of the automaton \mathcal{A} is encoded as a pile of encryptions (where each key encodes a letter of the stack). Then, upon receiving such a pile of encryptions encrypted by q on channel c_a , the process $P_{\mathcal{A}}$ will mimic the transition of \mathcal{A} that is triggered when the automaton is at state q upon reading a with the stack corresponding to that pile of encryptions.

More formally, the process $P_{\mathcal{A}}$ is defined as follows:

$$\begin{aligned}
 P_{\mathcal{A}} &\stackrel{\text{def}}{=} ! \text{in}(c_0, \text{start}).\text{new } \tilde{r}.\text{out}(c_0, \text{senc}(\text{senc}(\text{start}, k_{\omega}, r_1), k_{\text{start}}, r_2), q_0, r_3)) \quad (0) \\
 &\quad | ! \text{in}(c_a, \text{senc}(\overline{x.u}, q, z)).\text{new } \tilde{r}.\text{out}(c_a, \text{senc}(\overline{x.v}, q', r)) \quad (1) \\
 &\quad | ! \text{in}(c_a, \text{senc}(\overline{x.u'}, q, z)).\text{new } r.\text{out}(c_a, \text{senc}(\text{start}, k_{\text{well}}, r)) \quad (1a) \\
 &\quad | ! \text{in}(c_a, \text{senc}(\text{start}, k_{\text{well}}, z)).\text{new } r.\text{out}(c_a, \text{senc}(\text{start}, k_{\text{well}}, r)) \quad (1b) \\
 &\quad | ! \text{in}(c_f, \text{senc}(x, q_f, z)).\text{new } r.\text{out}(c_f, \text{senc}(\text{start}, q_f, r)), \quad (2)
 \end{aligned}$$

where a quantifies over Π , q over \mathcal{Q} , u over words in Γ^* such that $(q, a, u) \in \text{dom}(\delta)$, q_f over \mathcal{Q}_f , and $(q', v) = \delta(q, a, u)$. Lastly, u' ranges over $U'_{q,a} \stackrel{\text{def}}{=} \alpha \cdot SS_{q,a} \setminus S_{q,a}$, where $S_{q,a}$ ($SS_{q,a}$, respectively) is the set that contains suffixes (strict suffixes, respectively) of some u with $(q, a, u) \in \text{dom}(\delta)$. This set $U'_{q,a}$ corresponds intuitively to the set of shortest words that are not suffixes of any word in $\{u \mid (q, a, u) \in \text{dom}(\delta)\}$ and, thus, the shortest words to unstack to be sure that no transition from q reading a is possible in the automaton.

Example 6.1. Consider a real-time GPDA such that $\Gamma = \{\alpha, \beta, \gamma, \omega\}$, $q \in \mathcal{Q}$, and $a \in \Pi$. Assume that $\{u \mid (q, a, u) \in \text{dom}(\delta)\} = \{\beta\alpha, \beta\alpha\alpha\}$. We have $SS_{q,a} = \{\epsilon, \alpha, \alpha\alpha\}$ and $S_{q,a} = SS_{q,a} \cup \{\beta\alpha, \beta\alpha\alpha\}$. Thus, we have that

$$U'_{q,a} = \{\omega, \beta, \gamma, \omega\alpha, \gamma\alpha, \omega\alpha\alpha, \alpha\alpha\alpha, \gamma\alpha\alpha\}.$$

In the previous encoding, the branches (0) and (1) mimic the behavior of the automaton \mathcal{A} . Branch (2) is triggered in case a final state q_f is reached. In case we are considering a behavior that is not authorized by the automaton, we obtain a message encrypted with k_{well} through branches (1a). Then branches (1b) allow us to pursue the execution of the protocol outputting messages that look fresh.

LEMMA 6.2. *The protocol $P_{\mathcal{A}}$ described earlier is in \mathcal{C}_{pp} and of size polynomial w.r.t. \mathcal{A} .*

PROOF. First, note that because $\text{dom}(\delta)$ is finite, as the automaton is finitely described, the sets $\{u \mid (q, a, u) \in \text{dom}(\delta)\}$ and $U'_{q,a}$ are also finite for any $a \in \Pi$ and $q \in \mathcal{Q}$. Moreover, the automaton being deterministic, given $q \in \mathcal{Q}$ and $a \in \Pi$, for every word $s \in \Gamma^*$:

- either there exists a unique suffix u of s such that $(q, a, u) \in \text{dom}(\delta)$;
- or there exists a unique suffix u' of s such that $u' \in U'_{q,a}$,

and this disjunction is exclusive. This allows us to ensure that condition (2) of Definition 3.1 is satisfied, and thus $P_{\mathcal{A}}$ belongs to \mathcal{C}_{pp} .

Regarding the size of the protocol, the only nontrivial point is to check that the number of branches (1a) is polynomially bounded. Let $q \in \mathcal{Q}$ and $a \in \Pi$, and assume that the maximal length of a word u in a transition $q \xrightarrow{a;u/v} q'$ of the automaton is $\ell_{q,a}$; we have that the number of branches (1a) for state q and letter a is bounded by $\ell_{q,a} \times \#\Gamma \times \#\{u \mid (q, a, u) \in \text{dom}(\delta)\}$, where $\#S$ is the cardinality of set S . This allows us to conclude. \square

This polynomial encoding preserves inclusion.

PROPOSITION 6.3. *Let \mathcal{A} and \mathcal{B} be two real-time GPDA's. We have that*

$$\mathcal{L}(\mathcal{A}) \subseteq \mathcal{L}(\mathcal{B}) \iff P_{\mathcal{A}} \sqsubseteq P_{\mathcal{B}}.$$

PROOF. Let $\mathcal{A} = (Q, \Pi, \Gamma, q_0, \omega, Q_f, \delta)$ and $\mathcal{B} = (Q', \Pi, \Gamma', q'_0, \omega, Q'_f, \delta')$. We show the two implications separately.

(\Leftarrow) Assume that there exists a word $t \in \mathcal{L}(\mathcal{A}) \setminus \mathcal{L}(\mathcal{B})$. We will build a trace $(\text{tr}, \phi) \in \text{trace}(P_{\mathcal{A}})$ such that there exists no trace $(\text{tr}, \psi) \in \text{trace}(P_{\mathcal{B}})$, allowing us to conclude that $P_{\mathcal{A}} \not\sqsubseteq P_{\mathcal{B}}$. To build (tr, ϕ) , we will mimic the behavior of \mathcal{A} when reading t . The first branch to use is (0), enabling the attacker to activate other branches of the process $P_{\mathcal{A}}$. As $t \in \mathcal{L}(\mathcal{A})$ and \mathcal{A} is deterministic, there exists a unique sequence of transitions leading to an accepting state $q_f \in Q_f$. For every such transition the attacker will activate the corresponding branch (1) in $P_{\mathcal{A}}$. If $t = a_1 \dots a_n$, we define (tr, ϕ) as follows:

$$\text{tr} = \text{io}(c_0, \text{start}, w_1). \text{io}(c_{a_1}, w_1, w_2) \dots \text{io}(c_{a_n}, w_n, w_{n+1}). \text{io}(c_f, w_{n+1}, w_{n+2}),$$

and ϕ is defined as expected given our semantics. Because of the definition of the branch (1), the inputs on the channels c_{a_i} are possible, the stack of the automaton upon reading a_i and its current state being faithfully represented by the term $w_i\phi$. Thus, (tr, ϕ) is indeed a trace of $P_{\mathcal{A}}$. When reaching q_f , the attacker can use the branch (2) and output the message $\text{senc}(\text{start}, q_f, r)$. As $t \notin \mathcal{L}(\mathcal{B})$, the corresponding sequence of transitions in \mathcal{B} does not lead to any accepting state:

- either at some point of the execution of the automaton a transition from state q reading a is not possible with the current stack s . This means that there does not exist a suffix u of s such that $(q, a, u) \in \text{dom}(\delta')$, and thus, by definition of $U'_{q,a}$, there exists a suffix u' of s such that $u' \in U'_{q,a}$, enabling a transition (1a) on channel c_a for the attacker, and every subsequent transition is done using branches (1b);
- or the state reached in \mathcal{B} after reading t is not an accepting state, that is, not in Q'_f : the sequence $\text{in}(c_f, w_{n+1}). \text{out}(c_f, w_{n+2})$ cannot occur in $P_{\mathcal{B}}$.

Consequently, there exists no trace $(\text{tr}, \psi) \in \text{trace}(P_{\mathcal{B}})$ (for any ψ), and thus $P_{\mathcal{A}} \not\sqsubseteq P_{\mathcal{B}}$.

(\Rightarrow) First, note that for every frame ϕ (ψ , respectively) such that $(\text{tr}, \phi) \in \text{trace}(P_{\mathcal{A}})$ ($(\text{tr}, \psi) \in \text{trace}(P_{\mathcal{B}})$, respectively), we have that ϕ (ψ , respectively) is of the form

$$\{w_1 \triangleright \text{senc}(m_1, k_1, r_1), \dots, w_n \triangleright \text{senc}(m_n, k_n, r_n)\},$$

where the k_i are nondeducible and the r_i are “fresh” in the sense that they are all distinct and nondeducible. This means that no equality (but the trivial ones) holds in such a frame. Now consider the shortest trace $(\text{tr}, \phi) \in \text{trace}(P_{\mathcal{A}})$, in terms of number of transitions, such that there exists no equivalent frame $(\text{tr}, \psi) \in \text{trace}(P_{\mathcal{B}})$. Since keys are nondeducible, we may assume w.l.o.g that $(\text{tr}, \phi) \in \text{trace}_{\text{fwd}}(P_{\mathcal{A}})$. Because of the branches (1), (1a), and (1b) and in particular the definition of $U'_{q,a}$, for any $q \in Q$ and for any $a \in \Pi$, a transition of channel c_a is always possible, and we have seen that the resulting frames are necessarily in static equivalence. Thus, only the shortest trace where $P_{\mathcal{B}}$ will not be able to follow is when tr ends with an input/output on channel c_f . Let $w \in \text{dom}(\phi)$ be the corresponding variable in the frame ϕ . Consider the subsequence $\text{seq}_{\text{tr}}(w)$ of tr and more precisely the sequence of channels that occurs in this subsequence. Such a sequence is of the form $c_0.c_{a_1} \dots c_{a_n}.c_f$.

Let $v = a_1 \dots a_n$. We have that v is a word of Π^* , and, in particular,

- $v \in \mathcal{L}(\mathcal{A})$: indeed, branches (1) in $P_{\mathcal{A}}$ faithfully represent transitions $(q, a, u) \in \text{dom}(\delta)$ and a branch (2) can only be fired if $q_f \in Q_f$;

— $v \notin \mathcal{L}(\mathcal{B})$: indeed, branch (2) could not be fired, and either \mathcal{B} cannot read v or, after reading v , \mathcal{B} is not in any state of \mathcal{Q}_f .

Hence, $v \in \mathcal{L}(\mathcal{A}) \setminus \mathcal{L}(\mathcal{B})$, proving that $\mathcal{L}(\mathcal{A}) \not\subseteq \mathcal{L}(\mathcal{B})$. \square

Therefore, checking for equivalence of protocols is as difficult as checking equivalence of real-time generalized pushdown deterministic automata.

7. IMPLEMENTATION

In this section, we detail our tool Cpp2dpa to convert protocols in \mathcal{C}_{pp} into GDPA, available online at

<http://www.lsv.ens-cachan.fr/~chretien/cpp2dpa.php>.

This tool takes two protocols in \mathcal{C}_{pp} as input, turns them into GDPAs, and, through the tool 1A1B1C [Henry and Sénizergues 2013], outputs whether the two protocols were in equivalence, yielding a witness of nonequivalence in the negative case in the form of a sequence of channels leading to an attack. The tool focuses on the encoding as described in Section 5. In particular, we assume the prior steps of Section 4 were successfully applied to the pair of protocols; namely, the bijection α as in Lemma 4.12 was successfully guessed and the oracles of Section 4.2.2 correctly added.

The tool Cpp2dpa is written in Python 3. From pairs of protocols in \mathcal{C}_{pp} , it generates three pairs of *normalized deterministic pushdown automata*, instead of directly two pairs of GDPAs (as described in Section 5). This was necessary to interface with 1A1B1C and involves no loss of generality, as the former are more expressive than our GDPA. The normalization process still has the inconvenience, in order to preserve the determinacy of the result, to output automata that may duplicate actions. More specifically, when necessary, the channels appearing in the potential witness of nonequivalence may be doubled. This technical detail does not impair the ability for the combined tool to prove equivalence or find witnesses, nevertheless.

7.1. Encoding Pairs

Most protocols use pairs. While our formalism does not directly support pairs, we may encode a restricted kind of pairing, when there are only constants (such as identities) on the right. Formally, this amounts to encoding a pair $\langle t, a \rangle$, where t is a term and a some constant, by an encryption $\text{senc}(t, a, r)$ for some random seed r . Provided constants used in concatenation are disjoint from constants used as keys, this encoding does not introduce any confusion. Note that since encryption is randomized, this pairing operator also differs as it is randomized.

7.2. Biometric Passport

We are interested here in proving the unlinkability of the electronic passport protocol. A detailed specification of it can be found in Arapinis et al. [2010]. Here, we only consider the passport's role and forget about the reader. The first case we consider is the flawed version corresponding to the French implementation of the passport, in which an attack arises from the ability of the attacker to observe whether a MAC check succeeds or not. As our framework does not directly enable us to deal with pairs of messages with their MAC, we model it by a signature: the attacker is able to obtain the plaintext of it (which amounts to retrieving the first component of the real pair) but cannot forge it (the attacker is not a priori able to forge a valid MAC). The resulting

process is defined as follows:

$$\begin{aligned}
 P_A &\stackrel{\text{def}}{=} ! \text{in}(c_1, \text{start}).\text{new } \tilde{r}'. \\
 &\quad \text{out}(c_1, \text{sign}(\text{senc}(\text{senc}(\text{senc}(n_r, k_r, r'_1), n_p, r'_2), k_E, r'_3), \text{mac}_{k_m}, r'_4)) \quad (1) \\
 &\quad | ! \text{in}(c_2, \text{sign}(\text{senc}(x, k_E, z_1), \text{mac}_{k_m}, z_2)).\text{new } r_5.\text{out}(c_2, \text{sign}(x, \text{mac}_{\text{ok}}, r_5)) \quad (2a) \\
 &\quad | ! \text{in}(c'_2, \text{sign}(\text{senc}(x, n_p, z_1), \text{mac}_{\text{ok}}, z_2)).\text{new } \tilde{r}''. \\
 &\quad \quad \text{out}(c'_2, \text{sign}(\text{senc}(\text{senc}(x, n_p, r''_1), k_p, r''_2), \text{mac}_{k_m}, r''_3)), \quad (2b)
 \end{aligned}$$

where $\text{new } \tilde{r}$ is a shortcut of $\text{new } r_1.\text{new } r_2.\text{new } r_3.\text{new } r_4$ (and similarly for $\text{new } \tilde{r}'$ and $\text{new } \tilde{r}''$). The protocol is modeled through three rules. Branch (1) corresponds to a message from the current session, emitted by the reader. While the original protocol can check the authenticity of the MAC and the value of the nonce sent to the passport, our formalism requires us to separate this into two steps: branches (2a) and (2b). Branch (2a) checks the validity of the MAC: if it is valid, it sends a message signed with mac_{ok} . On the other hand, branch (2b) checks the value of the nonce (i.e., n_p) and finally emits the last message of the protocol. To retrieve the attack, we introduce the message sent by the reader from a previous session with a new branch denoted (0):

$$\begin{aligned}
 & ! \text{in}(c_0, \text{start}).\text{new } \tilde{r}. \\
 & \quad \text{out}(c_0, \text{sign}(\text{senc}(\text{senc}(\text{senc}(n_r^0, k_r^0, r_1), n_p^0, r_2), k_E, r_3), \text{mac}_{k_m}, r_4)). \quad (0)
 \end{aligned}$$

Another protocol P_B is obtained by replacing mac_{k_m} by $\text{mac}_{k'_m}$ in branches (1), (2a), and (2b). Our tool Cpp2dpa can automatically check that $P_A \not\approx P_B$.

Another version P'_A is obtained by replacing branches (2a) and (2b) by the branch

$$\begin{aligned}
 & ! \text{in}(c_2, \text{sign}(\text{senc}(\text{senc}(x, n_p, z_1), k_E, z_2), \text{mac}_{k_m}, z_3)).\text{new } \tilde{r}''. \\
 & \quad \text{out}(c_2, \text{sign}(\text{senc}(\text{senc}(x, n_p, r''_1), k_E, r''_2), \text{mac}_{k_m}, r''_3)). \quad (2)
 \end{aligned}$$

The protocol P'_B is similarly defined, with $\text{mac}_{k'_m}$ instead of mac_{k_m} , in branch (2). This version models the safe implementation of the protocol, where the success or failure of the MAC check is invisible to the attacker. Our tool Cpp2dpa can automatically prove that $P'_A \approx P'_B$.

7.3. Experiments

We have tested our tool Cpp2dpa on the running example as defined in Example 2.7 and Example 4.5, as well as on an encoding of the electronic passport protocol, described in Section 7.2 in two versions, unsafe and safe (see Arapinis et al. [2010] for more details).

	Automata (in ms)	Grammars (in s)	Equivalence (in s)
Example 2.7	7.1	9.2	3,462 (attack)
Example 4.5	7.0	3.1	9,788 (proof)
Unsafe passport	7.1	23.2	4.89 (attack)
Safe passport	8.1	15.0	76.1 (proof)

The experiments were conducted on an Intel(R) Xeon(R) CPU X5650 @ 2.67GHz with 47 Go of RAM, using one core only. The first column corresponds to the cumuled time required to produce the different automata, the second one the time needed to convert the automata into grammars to be processed by 1A1B1C, and the third one to the status of the equivalence (proof or witness of nonequivalence) and the cumuled time spent to prove the equivalence (when it is the case) or the execution time to find a witness of nonequivalence, when possible. There is nonequivalence as soon as one of our three pairs of automata are not in equivalence. Since we execute 1A1B1C in parallel for each

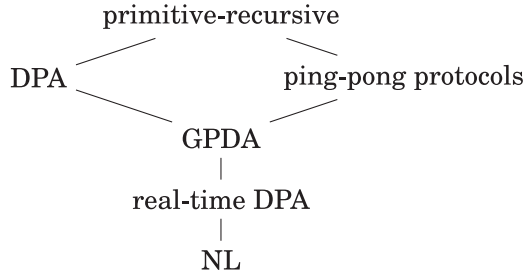


Fig. 6. Complexity bounds for equivalence of ping-pong protocols.

of these three pairs, the execution time corresponds to the first pair that is found to be not in equivalence. Converting the automata to grammars required an optimization of the built-in functionality in 1A1B1C in order to reach reasonable execution times. Other protocols were considered, namely, variants of the Wide Mouthed Frog, Denning-Sacco, and Private Authentication protocols. Unfortunately, for these ones, although the generation of automata was quick, it was impossible to prove (non)equivalence in reasonable time with 1A1B1C.

8. DISCUSSION AND CONCLUSION

We have shown a first decidability result for equivalence of (deterministic) ping-pong protocols for an unbounded number of sessions by reducing it to the equality of languages of deterministic, generalized, real-time pushdown automata (GPDA). We further showed that deciding equivalence of ping-pong protocols is actually at least as hard as deciding equality of languages of GPDA. Complexity-wise, the situation is slightly less clear. While the reduction from GPDA to ping-pong protocols is polynomial, the reduction from ping-pong protocols to GPDA requires an exponential blow-up. Indeed, to get rid of the attacker, we guess a correspondence between the keys of P and Q , and exponentially many such correspondences should be checked. In addition, the complexity of equivalence of various classes of pushdown automata is not very well known. It follows that the exact complexity of checking equivalence of protocols is unknown. The only upper bound is that equivalence is at most primitive recursive. This bound comes from the algorithm proposed by C. Stirling for equivalence of DPA [Stirling 2002]. The lower bound comes from the fact that real-time deterministic pushdown automata are at least NL-hard [Boehm and Goeller 2011]. Whether equivalence of DPA (or even real-time GPDA) is, for example, at least NP-hard is unknown. The complexity hierarchy known so far for equivalence of ping-pong protocols is displayed in Figure 6

Note that the complexity of GPDA and ping-pong protocols is actually quite close since the reduction from ping-pong protocols to GPDA is “just” exponential. Moreover, assume now that we consider only procedures that return a witness of nonequivalence (if any). Then the complexity classes of GPDA and ping-pong protocols should actually coincide. Indeed, assume that there is a procedure for checking equivalence of GPDA that ends in time $f(n)$, where n is the size of the inputs, and that returns a witness when two automata are not in equivalence. This witness must be of size at most $f(n)$. Then, given two ping-pong protocols P and Q , we would construct \bar{P} and \bar{Q} as defined in Lemma 4.12 step by step.

Instead of guessing the sets K and K' , we would start from the emptysets $K = K' = \emptyset$. If $\bar{P} \not\approx \bar{Q}$, that is, if $\mathcal{A}_{\bar{P}} \not\approx \mathcal{A}_{\bar{Q}}$, we consider a witness of nonequivalence. Either it is a witness of $P \not\approx Q$ (and we are done) or there must exist a key k that is deducible in P and a corresponding key k' deducible with the same actions in Q . We start over with $K = \{k\}$ and $K' = \{k'\}$.

This algorithm has at most n steps, and each step involves a call to the GPDA procedure ($\mathcal{A}_{\bar{P}} \approx \mathcal{A}_{\bar{Q}}$) and involves replaying a witness of size $f(n)$. This yields a procedure of complexity $O(f(n))$.

Our class of security protocols handles only randomized primitives, namely, symmetric/asymmetric encryptions and signatures. Our decidability result could be extended to handle deterministic primitives instead of the randomized one (the reverse encoding—from real-time GPDAs to processes with deterministic encryption—may not hold anymore). Due to the use of pushdown automata, extending our decidability result to protocols with pairing is not straightforward. A direction is to use pushdown automata for which stacks are terms.

While we consider an unbounded number of sessions, we consider a fixed number of agents in our examples. We could model an unbounded number of agents; however, since our class considers protocol rules with at most one variable, we could consider at most one agent per rule with no key nor nonces, which would be very restrictive. Another direction is to study whether we can soundly bound the number of agents.

Our tool Cpp2dpa in combination with 1A1B1C yields the first implementation of a decidability procedure for equivalence of protocols for an unbounded number of sessions. However, the number of protocols covered so far is limited. A first reason is in the limitations of the class of ping-pong protocols. However, another reason is the (too long) time needed to check for equivalence. Our transformation from protocols to automata using Cpp2dpa remains reasonably fast. Most of the execution time comes from 1A1B1C. Since this tool is still in its early stage of development, we may hope for significant improvement of 1A1B1C's performance in the next years.

APPENDICES

A. UNDECIDABILITY OF TRACE INCLUSION

The purpose of this section is to establish the following result.

THEOREM 3.6. *The following problem is undecidable:*

Input. P and Q two protocols in \mathcal{C}_{pp} .

Output. Whether P is trace included in Q , that is, $P \sqsubseteq Q$.

An instance of the PCP over the alphabet A is given by two sets of tiles $U = \{u_i \mid 1 \leq i \leq n\}$ and $V = \{v_i \mid 1 \leq i \leq n\}$, where $u_i, v_i \in A^*$. The problem consists of deciding whether there exists a nonempty sequence i_1, \dots, i_p over $\{1, \dots, n\}$ such that $u_{i_1} \dots u_{i_p} = v_{i_1} \dots v_{i_p}$.

To prove the undecidability of trace inclusion in \mathcal{C}_{pp} , we show it is possible to encode the Post Correspondence Problem into an inclusion of two protocols of this class. Given a word, one protocol will be meant to unstack the first set of tiles, while the other will try as much as possible to unstack the second set of tiles. While an empty word is not “simultaneously” reached by the two processes, their traces appear to be equivalent. Conversely, if a solution to the Post Correspondence Problem does exist, it will lead the second process to react in a distinct way (by stopping its execution), breaking the trace inclusion property.

For each $i \in \{1, \dots, n\}$, we define two (possibly empty) sets of words over A , namely, $W_i \stackrel{\text{def}}{=} A^{|v_i|} \setminus \{v_i\}$ and $W'_i \stackrel{\text{def}}{=} A^0 \cup A^1 \cup \dots \cup A^{|v_i|-1}$, where $|v_i|$ denotes the length of the word v_i .

Example A.1. Let $A = \{a, b\}$ and consider the following pairs of tiles: (b, ϵ) , (b, a) , and (a, ba) . This instance of PCP admits a solution. Indeed, the nonempty sequence 13 leads to the word $u_1 u_3 = v_1 v_3 = ba$. We have $W_1 = W'_1 = \emptyset$, $W_2 = \{b\}$ and $W'_2 = \{\epsilon\}$, and lastly $W_3 = \{aa, ab, bb\}$ and $W'_3 = \{a, b, \epsilon\}$.

Words in A^* will be represented through nested symmetric encryption with private keys representing their counterparts in A . For the sake of brevity, given a word $u = \alpha_1 \dots \alpha_p$ of A^* , we denote by

- \bar{u} the term $\text{senc}(\dots \text{senc}(\epsilon, \alpha_1, z_1) \dots, \alpha_p, z_p)$; and
- $\overline{x.u}$ the term $\text{senc}(\dots \text{senc}(x, \alpha_1, z_1) \dots, \alpha_p, z_p)$,

where z_1, \dots, z_p are variables of sort rand . Note that if $u = \epsilon$, then $\bar{u} = \epsilon$, and $\overline{x.u} = x$.

Next, k_i, k'_i with $i \in \{0, 1, 2, 3\}$ are constants in Σ_0 of sort SymKey , and for each $\alpha \in A$, we denote also by α its counterpart in Σ_0 (constants of sort SymKey). We denote by ϵ a constant in Σ_0 of sort msg . These constants are initially unknown by the attacker, and actually it is quite easy to see that they will be never revealed. Lastly, c, c_α, c_i, c' with $\alpha \in A$ and $i \in \{1, \dots, n\}$ are constant symbols of sort channel in Ch .

Let P_U and P_V be the following protocols:

$$\begin{aligned}
 P_U := & \quad !\text{in}(c, \text{start}).\text{new } r.\text{out}(c, \text{senc}(\epsilon, k_0, r)) & (\text{start}) \\
 & | !\text{in}(c_\alpha, \text{senc}(x, k_0, z)).\text{new } r_1, r_2.\text{out}(c_\alpha, \text{senc}(\text{senc}(x, \alpha, r_2), k_0, r_1)) & (1) \\
 & | !\text{in}(c_i, \text{senc}(\overline{x.u_i}, k_0, z)).\text{new } r.\text{out}(c_i, \text{senc}(x, k_1, r)) & (2) \\
 & | !\text{in}(c_i, \text{senc}(\overline{x.u_i}, k_1, z)).\text{new } r.\text{out}(c_i, \text{senc}(x, k_1, r)) & (3) \\
 & | !\text{in}(c', \text{senc}(\epsilon, k_1, z)).\text{new } r.\text{out}(c', \text{senc}(\epsilon, k_2, r)), & (4)
 \end{aligned}$$

where i ranges in $\{1, \dots, n\}$ and α in A .

The branch (start) is the only way to start an execution, and then branches (1) are used to build a word $\alpha_1 \dots \alpha_n$ (that could be a Post word in case we consider a positive instance of PCP). This word will be represented through the term $\text{senc}(\dots \text{senc}(\epsilon, \alpha_1, r_1) \dots, \alpha_n, r_n)$ up to the choice of randoms. Then, branches (2) and (3) are used to unstack the different tiles u_1, \dots, u_n . Note that the purpose of having two similar branches (but using different keys) for this task is to ensure that we will unstack at least one tile, and thus the sequence $i_1 \dots i_p$ of indices is not empty. Then, reaching the empty word when unstacking these tiles will allow us to perform input/output on channel c' (branch (4)):

$$\begin{aligned}
 P_V := & \quad !\text{in}(c, \text{start}).\text{new } r.\text{out}(c, \text{senc}(\epsilon, k'_0, r)) & (\text{start}) \\
 & | !\text{in}(c_\alpha, \text{senc}(x, k'_0, z)).\text{new } r_1, r_2.\text{out}(c_\alpha, \text{senc}(\text{senc}(x, \alpha, r_2), k'_0, r_1)) & (1) \\
 & | !\text{in}(c_i, \text{senc}(\overline{x.v_i}, k'_0, z)).\text{new } r.\text{out}(c_i, \text{senc}(x, k'_1, r)) & (2') \\
 & | !\text{in}(c_i, \text{senc}(\overline{x.w}, k'_0, z)).\text{new } r.\text{out}(c_i, \text{senc}(\epsilon, k'_3, r)) & (2'a) \\
 & | !\text{in}(c_i, \text{senc}(\overline{w'}, k'_0, z)).\text{new } r.\text{out}(c_i, \text{senc}(\epsilon, k'_3, r)) & (2'b) \\
 & | !\text{in}(c_i, \text{senc}(\overline{x.v_i}, k'_1, z)).\text{new } r.\text{out}(c_i, \text{senc}(x, k'_1, r)) & (3') \\
 & | !\text{in}(c_i, \text{senc}(\overline{x.w}, k'_1, z)).\text{new } r.\text{out}(c_i, \text{senc}(\epsilon, k'_3, r)) & (3'a) \\
 & | !\text{in}(c_i, \text{senc}(\overline{w'}, k'_1, z)).\text{new } r.\text{out}(c_i, \text{senc}(\epsilon, k'_3, r)) & (3'b) \\
 & | !\text{in}(c', \text{senc}(\overline{x.\beta}, k'_1, z)).\text{new } r.\text{out}(c', \text{senc}(\epsilon, k'_2, r)) & (4'a) \\
 & | !\text{in}(c', \text{senc}(\epsilon, k'_3, z)).\text{new } r.\text{out}(c', \text{senc}(\epsilon, k'_2, r)) & (4'b) \\
 & | !\text{in}(c_i, \text{senc}(\epsilon, k'_3, z)).\text{new } r.\text{out}(c_i, \text{senc}(\epsilon, k'_3, r)), & (5')
 \end{aligned}$$

where i ranges in $\{1, \dots, n\}$, α and β in A , and for each $i \in \{1, \dots, n\}$, w in W_i and w' in W'_i .

The protocol P_V has the same structure as P_U . However, it is more complex since we want P_V to follow the execution of P_U as soon as the execution does not correspond to a solution of the PCP problem. In particular, we do not want P_V to block in case it is not able to unstack a tile v_i . To achieve this, some additional branches are added (namely, (2'a) and (2'b), as well as (3'a) and (3'b)). Intuitively, those branches are triggered when

(2') and (3') cannot be, and the resulting term is encrypted with a special key k'_3 that will allow P_V to mimic the remainder of the execution using branch (5'). Now, regarding the branches on channel c' , the idea is to allow P_V to mimic the behavior of P_U only when the trace tr does not correspond to a solution of the PCP. To achieve this, we allow P_V to follow P_U only when the term given in input on channel c' is not a legal encoding of the empty word. Such a term will go through (4'a) or (4'b).

Note that, on both protocols, the terms that are outputted look like fresh random numbers due to fresh nonces occurring in every output and ignorance of the keys. In other words, the two frames resulting from the execution of, respectively, P_U and P_V always remain in static equivalence. Therefore, checking trace equivalence amounts to checking that any execution trace of P_U is a trace of P_V , and conversely.

LEMMA A.2. *The protocols P_U and P_V described previously are in \mathcal{C}_{pp} .*

PROOF. The only nontrivial point is to ensure that condition (2) stated in Definition 3.1 is satisfied, that is, to ensure that pattern matching operated by inputs taking place on the same channel is exclusive. Regarding protocol P_U , when two inputs occur on the same channel c_i , we have that the outermost key is different. Regarding protocol P_V , the result also holds thanks to the exclusivity of the pattern matching obtained through a careful definition of sets W_i and W'_i . For instance, note that when $v_i = \epsilon$, $W_i = W'_i = \emptyset$, and thus there is no branch (2'a)/(2'b) ((3'a)/(3'b), respectively). \square

PROPOSITION A.3. *Let U/V be an instance of PCP. We have that $P_U \sqsubseteq P_V$ if and only if U/V is a negative instance of PCP (i.e., an instance with no solution).*

PROOF. We prove successively the two implications.

(\Rightarrow) *If U/V is a positive instance of PCP, then $P_U \not\sqsubseteq P_V$.* If U/V is a positive instance of PCP, there exists a nonempty sequence $i_1 \dots i_p$ over $\{1, \dots, n\}$ such that $u_{i_1} \dots u_{i_p} = v_{i_1} \dots v_{i_p}$.

Let $u = \alpha_1 \dots \alpha_m$ be the resulting word over A . From this word and the sequence i_1, \dots, i_p , the attacker playing with P_U can build the term $\text{senc}(\bar{u}, k_0, r)$ representing the word u with branches (1) and then remove one by one the tiles u_{i_p} to u_{i_1} using (2) and (3). Let tr be the resulting trace of the protocol P_U :

$$\begin{aligned} \text{tr} \stackrel{\text{def}}{=} & \text{io}(c, \text{start}, w_1). \text{io}(c_{\alpha_1}, w_1, w_2) \dots \text{io}(c_{\alpha_m}, w_m, w_{m+1}) \\ & \text{io}(c_{i_p}, w_{m+1}, w_{m+2}) \dots \text{io}(c_{i_1}, w_{p+m}, w_{p+m+1}). \text{in}(c', w_{m+p+1}), \end{aligned}$$

where $\text{io}(c, R, w) \stackrel{\text{def}}{=} \text{in}(c, R). \text{out}(c, w)$.

The trace tr models the fact that, given $\text{senc}(\bar{u}, k_0, r)$ (stored in w_{m+1}), P_U can remove one by one the tiles u_{i_p} to u_{i_1} to reach the empty word and hence output the message $\text{senc}(\epsilon, k_1, r)$ (stored in w_{m+p+1}) that can then be accepted as input on c' . In this execution, no equality holds in the resulting frame ϕ , as the attacker ignores the keys that are used to encrypt, and all outputted messages use different random seeds; thus, all messages look fresh.

We claim that this trace does exist in P_V ; that is, there exists no ψ such that $(\text{tr}, \psi) \in \text{trace}(P_V)$. Indeed, the pattern matching operated by P_V is exclusive once the term and the channel are fixed. Thus, P_V has no choice but to remove tiles v_{i_p} to v_{i_1} using (2') and (3') leading to the term $\text{senc}(\epsilon, k'_1, r)$ (stored in w_{m+p+1}) as $\alpha_1 \dots \alpha_m$ is a Post word. Any other trace would lead to either a mismatch on the channels or an improper filtering in P_V . Then the action $\text{in}(c', w_{m+p+1})$ will have no counterpart on P_V . So (tr, ϕ) has no equivalent trace in P_V , that is, $P_U \not\sqsubseteq P_V$.

(\Leftarrow) *If U/V is a negative instance of PCP, then $P_U \sqsubseteq P_V$.* Let $(\text{tr}, \phi) \in \text{trace}(P_U)$; we aim at showing that there exists an equivalent trace $(\text{tr}, \psi) \in \text{trace}(P_V)$. Actually, since

terms that are outputted by P_U and P_V look like fresh random numbers, we simply have to show that there exists ψ such that $(\text{tr}, \psi) \in \text{trace}(P_V)$. Two cases can occur for any trace $(\text{tr}, \phi) \in \text{trace}(P_U)$:

- tr contains no input on channel c' . In such a case, by construction of P_V , the frame ψ can be built by following the sequence of channels used in tr and choosing the adequate filtering. It is always possible to do so, as the definition of sets W_i and W'_i ensure that every term built by the attacker can be handled on any channel c_i . Note that when the term given in input is of the form $\text{senc}(\epsilon, k'_3, r)$ for some r , it would be accepted on any channel.
- tr contains an input on channel c' . In such a case, this means that the associated term $\text{senc}(\overline{\alpha_1 \dots \alpha_m}, k_0, r)$ that has been built using channels c_α with $\alpha \in A$ is a word made of tiles in $\{u_1, \dots, u_n\}$. Indeed, the only way to activate an input on c' is to go through the branches (2) and (3) by unstacking the said tiles. Then, because this particular instance of PCP has no solution, such a word $\alpha_1 \dots \alpha_m$ cannot be a Post word and thus it cannot be decomposed using tiles in $\{v_1, \dots, v_n\}$ following the same sequence of indices: because the filtering in P_V is also exhaustive, messages outputted by P_V from a certain point will be either encrypted by k'_3 or will reach the end of the sequence with a term of the form $\text{senc}(u, k'_1, r)$ with u different from the constant ϵ . Thanks to branches (4'a), (4'b), and (5'), P_V will be able to follow P_U .

Hence, for any trace $(\text{tr}, \phi) \in \text{trace}(P_U)$, there exists a trace $(\text{tr}, \psi) \in \text{trace}(P_V)$. It remains to show that $\phi \sim \psi$. This is due to the fact that both ϕ and ψ are of the form $\{w_1 \triangleright \text{senc}(m_1, k_1, r_1), \dots, w_n \triangleright \text{senc}(m_n, k_n, r_n)\}$, where the k_i are nondeducible and the r_i are “fresh” in the sense that they are all distinct and nondeducible. We therefore conclude that $P_U \sqsubseteq P_V$. \square

Theorem 3.6 directly follows from Proposition A.3 and the undecidability of the Post Correspondence Problem.

B. GETTING RID OF THE ATTACKER

LEMMA B.1. *Let P and Q be two protocols in \mathcal{C}_{pp} and \mathcal{K}_P (\mathcal{K}_Q , respectively) be the set of deducible constants of sort key that occur in P (Q , respectively); if $P \approx Q$, then there exists a unique bijection α from \mathcal{K}_P to \mathcal{K}_Q such that for every trace $(\text{tr}, \phi) \in \text{trace}(P)$, there exists a trace $(\text{tr}, \psi) \in \text{trace}(Q)$ such that for any recipe R and any $k \in \mathcal{K}_P$:*

- $R\phi \downarrow$ is of sort s if and only if $R\psi \downarrow$ is of sort s ,
 - $R\phi \downarrow = k$ if and only if $R\psi \downarrow = \alpha(k)$; and
 - $R\phi \downarrow = k^{-1}$ if and only if $R\psi \downarrow = (\alpha(k))^{-1}$;
- where $s \in \{\text{SymKey}, \text{PubKey}, \text{PrivKey}\}$;*

and conversely, for every $(\text{tr}, \psi) \in \text{trace}(Q)$, there exists a trace $(\text{tr}, \phi) \in \text{trace}(P)$ satisfying the same properties.

PROOF. We can describe α as a relation in the following way:

for every $k \in \mathcal{K}_P$ of sort s , and every trace $(\text{tr}, \phi) \in \text{trace}(P)$ and recipe R such that $R\phi \downarrow = k$, we define $\alpha(k) = R\psi \downarrow$, where ψ is the only frame such that $(\text{tr}, \psi) \in \text{trace}(Q)$.

The existence of such a frame comes from the fact that $P \approx Q$, whereas its unicity is a consequence of the determinism of protocols in \mathcal{C}_{pp} .

We now need to prove that our definition of α is sound and unambiguous. To do so, we show that

- $R\psi\downarrow$ is a constant of sort s . We have that there exists a trace $(tr, \phi) \in \text{trace}(P)$ such that $R\phi\downarrow = k \in K_P$. Since $P \approx Q$ and Q is in \mathcal{C}_{pp} , we consider the trace $(tr, \psi) \in \text{trace}(Q)$. By definition of static equivalence, we have that $R\psi\downarrow$ is a constant of sort s . Otherwise, we would have that $\text{senc}(\text{start}, R, r_i)\phi \in \mathcal{T}(\Sigma, \mathcal{N})$, whereas $\text{senc}(\text{start}, R, r_i)\psi \notin \mathcal{T}(\Sigma, \mathcal{N})$ if $s = \text{SymKey}$ (the resulting term is not properly sorted). The same argument applies with aenc and sign for s equal to PubKey and PrivKey , respectively.
- We have that $|\mathcal{K}_P| = |\mathcal{K}_Q|$. Suppose *ad absurdum* that, for instance, $|\mathcal{K}_P| < |\mathcal{K}_Q|$. Since every element of \mathcal{K}_Q is deducible (and due to the shape of the protocols under study), there exists $(tr, \psi) \in \text{trace}(Q)$ such that for all $k \in \mathcal{K}_Q$, there exists a recipe R_k such that $R_k\psi\downarrow = k$. In particular, when $k \neq k'$, we have that $R_k\psi\downarrow \neq R_{k'}\psi\downarrow$. Since $P \approx Q$, there exists a frame ϕ such that $(tr, \phi) \in \text{trace}(P)$. Thanks to the previous item, we know that $R_k\phi\downarrow$ ($R_{k'}\phi\downarrow$, respectively) has the same sort as $R_k\psi\downarrow$ ($R_{k'}\psi\downarrow$, respectively), that is, sort key. As $|\mathcal{K}_P| < |\mathcal{K}_Q|$, there exist two distinct keys k and k' such that $R_k\phi\downarrow = R_{k'}\phi\downarrow$. Hence, ϕ and ψ are not statically equivalent, contradicting the fact that $P \approx Q$. The case where $|\mathcal{K}_Q| < |\mathcal{K}_P|$ can be handled similarly.
- α is a function. Suppose there exist a trace $(tr, \phi) \in \text{trace}(P)$, a recipe R_i , and a corresponding equivalence trace $(tr, \psi) \in \text{trace}(Q)$ such that $R_i\phi\downarrow = k$ and $R_i\psi\downarrow = k'$; and a trace $(tr', \phi') \in \text{trace}(P)$, a recipe R_j , and a corresponding equivalence trace $(tr', \psi') \in \text{trace}(Q)$ such that $R_j\phi'\downarrow = k$ but $R_j\psi'\downarrow = k''$ with $k' \neq k''$. Considering the trace made up of the trace tr followed by tr' , it is then possible to exhibit a witness of nonequivalence. More precisely, relying on R_i and R_j , we can build a test that holds in the resulting frame when executing P , whereas this test will not hold on the frame resulting from the execution of Q .

Now we show that α is an injection, that is, $\alpha(k) \neq \alpha(k')$ as soon as k, k' are two distinct elements of \mathcal{K}_P . Suppose, as previously, there exist a trace $(tr, \phi) \in \text{trace}(P)$, a recipe R_i , and a corresponding equivalence trace $(tr, \psi) \in \text{trace}(Q)$ such that $R_i\phi\downarrow = k$ and $R_i\psi\downarrow = \alpha(k)$; and a trace $(tr', \phi') \in \text{trace}(P)$, a recipe R_j , and a corresponding equivalence trace $(tr', \psi') \in \text{trace}(Q)$ such that $R_j\phi'\downarrow = k'$ but $R_j\psi'\downarrow = \alpha(k)$ with $k \neq k'$. Considering the trace made up of the trace tr followed by tr' , it is then possible to exhibit a witness of nonequivalence. More precisely, relying on R_i and R_j , we can build a test that holds in the frame resulting from the execution of P and that does not hold when executing Q . Thus, we have now proven that α is a bijection.

Note that we have already proved that: $R\phi\downarrow = k$ if and only if $R\psi\downarrow = \alpha(k)$.

To show that α satisfies the last condition (item 3), suppose that $k \in \mathcal{K}_P$, and $R\phi\downarrow = k^{-1}$. As previously shown, $R\psi\downarrow = \alpha(k^{-1})$. We want to prove that $\alpha(k^{-1}) = (\alpha(k))^{-1}$. If k is of sort SymKey , the result is obvious as $k^{-1} = k$ for any such key. Suppose k is of sort PubKey . We have now that there exists a trace $(tr, \phi) \in \text{trace}(P)$ and a recipe R' such that $R'\phi\downarrow = k \in K_P$. Since $P \approx Q$, consider the corresponding equivalence trace $(tr, \psi) \in \text{trace}(Q)$. Consider the recipes $R_1 = \text{aenc}(\text{start}, R', n)$ and $R_2 = \text{adec}(R_1, R)$. Then $R_2\phi\downarrow = \text{start}$ and $R_2\psi\downarrow = \text{start}$ if and only if $R\psi\downarrow = (R'\psi)^{-1}$. As we have already proved that α preserves sorts, we get that $R_2\psi\downarrow$ is of sort msg if and only if $\alpha(k^{-1}) = R\psi\downarrow = (R'\psi\downarrow)^{-1} = (\alpha(k))^{-1}$. Hence, α is compatible with the inverse function. The same argument can be used if k is of sort PrivKey with sign and check .

Finally, we prove the unicity of such a bijection: suppose there were α' an adequate bijection and $k \in \mathcal{K}_P$ such that $\alpha(k) \neq \alpha'(k)$. By definition of α , for every trace $(tr, \phi) \in \text{trace}(P)$ and every recipe R such that $R\phi\downarrow = k$, $\alpha(k) = R\psi\downarrow$. But as α' satisfies a similar property, we get that $R\psi\downarrow = \alpha'(k)$, contradicting our hypothesis. Hence, α is unique. Determinism of P and Q then ensures that traces of P and Q are uniquely matched (as $P \approx Q$), thus guaranteeing the converse part of the Lemma. \square

LEMMA B.2. *Let P and Q be two protocols in \mathcal{C}_{pp} respectively disclosing two sets of keys K and K' as in Lemma 4.12. Then $P \approx Q$ if and only if $\bar{P} \approx_{\text{fwd}} \bar{Q}$, where \bar{P} and \bar{Q} are as defined in Lemma 4.14. We call T_{oracle} the transformation taking a pair of protocols (P, Q) satisfying the aforementioned condition and returning the pair (\bar{P}, \bar{Q}) presently defined.*

PROOF. Let \mathcal{K}_P (\mathcal{K}_Q , respectively) be the set of deducible constants of sort key that occur in P (Q , respectively). We recall that, as a consequence of Lemma 4.12, we necessarily have that $\mathcal{K}_P \subseteq K$ and $\mathcal{K}_Q \subseteq K'$. Because protocols P and \bar{P} (Q and \bar{Q} , respectively) disclose all their deducible keys, there exists a trace (tr_0, ϕ_0) of P and \bar{P} $((\text{tr}_0, \psi_0)$ a trace of Q and \bar{Q} , respectively) defined as follows:

$$\text{tr}_0 = \text{in}(c_{k_1, \alpha(k_1)}, \text{start}).\text{out}(c_{k_1, \alpha(k_1)}, w_1^0) \dots \text{in}(c_{k_n, \alpha(k_n)}, \text{start}).\text{out}(c_{k_n, \alpha(k_n)}, w_n^0)$$

for $k_1, \dots, k_n \in \mathcal{K}_P$, and $\phi_0 = \{w_1^0 \triangleright k_1, \dots, w_n^0 \triangleright k_n\}$, and symmetrically for Q and \bar{Q} . In the following, we will assume that a trace of P or \bar{P} (of Q or \bar{Q} , respectively) starts with the prefix tr_0 and contains the frame ϕ_0 .

For sake of clarity of the construction explained later, we actually show that

$$\bar{P} \approx_{\text{fwd}} \bar{Q} \text{ if, and only if } P^+ \approx Q^+,$$

where $P^+ = P \mid \text{in}(c, x).\text{out}(c, x)$ and $Q^+ = Q \mid \text{in}(c, x).\text{out}(c, x)$ for some fresh channel name c . Then, it is easy to conclude the expected result relying on the fact that $P \approx Q$ is equivalent to $P^+ \approx Q^+$.

(\Rightarrow) First, suppose $\bar{P} \not\approx_{\text{fwd}} \bar{Q}$. Assume that there exists $(\text{tr}, \phi) \in \text{trace}_{\text{fwd}}(\bar{P})$ such that there is no equivalent frame ψ such that $(\text{tr}, \psi) \in \text{trace}_{\text{fwd}}(\bar{Q})$. We define $(\text{tr}', \phi) \in \text{trace}(P^+)$ as follows:

- every sequence $\text{in}(c_{k, \alpha(k)}^{\text{senc}}, R).\text{out}(c_{k, \alpha(k)}^{\text{senc}}, w')$ in tr is replaced by the sequence $\text{in}(c, \text{senc}(R, w_k^0, n)).\text{out}(c, w')$ in tr' , where n is a fresh name.
- every sequence $\text{in}(c_{k, \alpha(k)}^{\text{sdec}}, R).\text{out}(c_{k, \alpha(k)}^{\text{sdec}}, w')$ in tr is replaced by the sequence $\text{in}(c, \text{sdec}(R, w_k^0)).\text{out}(c, w')$ in tr' .
- every sequence $\text{in}(c_{k, \alpha(k)}^{\text{aenc}}, R).\text{out}(c_{k, \alpha(k)}^{\text{aenc}}, w')$ in tr is replaced by the sequence $\text{in}(c, \text{aenc}(R, w_k^0, n)).\text{out}(c, w')$ in tr' , where n is a fresh name.
- every sequence $\text{in}(c_{k, \alpha(k)}^{\text{adec}}, R).\text{out}(c_{k, \alpha(k)}^{\text{adec}}, w')$ in tr is replaced by the sequence $\text{in}(c, \text{adec}(R, w_k^0)).\text{out}(c, w')$ in tr' .
- every sequence $\text{in}(c_{k, \alpha(k)}^{\text{sign}}, R).\text{out}(c_{k, \alpha(k)}^{\text{sign}}, w')$ in tr is replaced by the sequence $\text{in}(c, \text{sign}(R, w_k^0, n)).\text{out}(c, w')$ in tr' , where n is a fresh name.
- every sequence $\text{in}(c_{k, \alpha(k)}^{\text{check}}, R).\text{out}(c_{k, \alpha(k)}^{\text{check}}, w')$ in tr is replaced by the sequence $\text{in}(c, \text{check}(R, w_k^0)).\text{out}(c, w')$ in tr' .

Note that by definition of a trace being in $\text{trace}_{\text{fwd}}(\bar{P})$, we have that R is either a variable w or the constant start . We claim that there exists no frame ψ such that $(\text{tr}', \psi) \in \text{trace}(Q^+)$ with $\phi \sim \psi$. Indeed, because the frames are left unchanged, the input recipes match the same input patterns, and recipes holding true and false keep their truth values. So if such a frame ψ existed, (tr, ψ) would belong to $\text{trace}_{\text{fwd}}(\bar{Q})$ and be equivalent to (tr, ϕ) .

(\Leftarrow) Now, suppose $P \not\approx Q$. We have that $P^+ \not\approx Q^+$, and we can even assume that $P^+ \not\approx^{\text{io}*} Q^+$. We consider a witness of this nonequivalence, that is, a trace tr such that $(\text{tr}, \phi) \in \text{trace}^{\text{io}*}(P^+)$ and for which there exists no equivalent frame ψ such that $(\text{tr}, \psi) \in \text{trace}^{\text{io}*}(Q^+)$. Actually, we can even assume w.l.o.g. that

- every input recipe in tr on a channel different from c is either a variable w or the constant start ;
- every input recipe in tr on channel c involves at most one function symbol in Σ_{pub} ; and
- $\phi \not\sim_{\text{fwd}} \psi$, that is, we consider recipes that are either variables or the constant start .

We consider the shortest trace $(\text{tr}, \phi) \in \text{trace}^{\text{io}*}(P)$, in terms of number of transitions, such that there is no equivalent frame ψ satisfying $(\text{tr}, \psi) \in \text{trace}^{\text{io}*}(Q)$, and for which all the requirements listed previously are satisfied.

Through recipes of the form $\text{senc}(u, v, w)$ on channel c , the attacker has the ability to use the same random seed more than once. Let us first show that we can always assume tr uses nonces at most once. If this is not the case, we build a new trace $(\tilde{\text{tr}}, \tilde{\phi})$, such that ϕ is statically equivalent to $\tilde{\phi}$ for which it is the case.

First, we consider the case where there exists no ψ such that $(\text{tr}, \psi) \in \text{trace}^{\text{io}*}(Q)$. Because random seeds are not filtered in protocols of \mathcal{C}_{pp} (every input pattern contains distinct variables as third argument), we can rename some occurrences of the random seeds of the attacker (i.e., the random seeds appearing in the recipes on channel c) by fresh random seeds without changing the status of the trace (i.e., the fact that the trace is executable or not). Given tr_ρ , such a trace obtained by renaming, we have that $(\text{tr}_\rho, \phi_\rho) \in \text{trace}^{\text{io}*}(P)$ for some frame ϕ_ρ whereas $(\text{tr}_\rho, \psi_\rho) \notin \text{trace}^{\text{io}*}(Q)$ for any frame ψ_ρ . And in particular, if we choose tr_ρ such that there are no two identical nonces in its image, we get a witness of nonequivalence with pairwise distinct random seeds for the attacker.

Now we consider the case where $(\text{tr}, \psi) \in \text{trace}^{\text{io}*}(Q)$ but $\phi \not\sim_{\text{fwd}} \psi$. Suppose r is a random seed that appears twice in tr in two contexts $f(w_i, w_j, r)$ and $f(w'_i, w'_j, r)$ for some $f \in \Sigma_{\text{pub}}$ with $w_i\phi = w'_i\phi$ and $w_j\phi = w'_j\phi$. Because tr is a minimal witness of nonequivalence, $\phi_{-1} \sim_{\text{fwd}} \psi_{-1}$, where ϕ_{-1} (ψ_{-1} , respectively) denotes ϕ (ψ , respectively) minus its last element. Consequently, we also have that $w_i\psi = w'_i\psi$ and $w_j\psi = w'_j\psi$, as $w_i, w_j, w'_i, w'_j \in \text{dom}(\phi_{-1})$ (they are used in input recipes). Let w and w' be the corresponding outputs of the recipes $f(w_i, w_j, r)$ and $f(w'_i, w'_j, r)$ and assume w appears before w' in tr : we now have that $w = w'$ in both ϕ and ψ , and we can safely replace any occurrence of w' in tr by w . The resulting trace is still a witness of nonequivalence as the substitution replaces identical terms in ψ .

Thus, it remains only to consider the case where a random seed appears twice in tr but such that either the function symbol, the plaintext, or the keys are different. Formally, consider the two contexts $f(w_i, w_j, r)$ and $g(w'_i, w'_j, r)$ with $f, g \in \Sigma_{\text{pub}}$, w and w' their respective output variables as before, and either $w_i\phi \neq w'_i\phi$, $w_j\phi \neq w'_j\phi$, or $f \neq g$. Following the same reasoning as before, as $\phi_{-1} \sim_{\text{fwd}} \psi_{-1}$, the same inequality has to hold in ψ . Consider the test $w_k = w'_k$, which distinguishes between ϕ and ψ : suppose $w_k\phi = w'_k\phi$ but $w_k\psi \neq w'_k\psi$. Replacing r by r' in $g(w'_i, w'_j, r)$ will still lead to $w_k\psi \neq w'_k\psi$ (after replacement) as no equality between subterms is added. But if $w_k\phi \neq w'_k\phi$ (after replacement), it would imply that there were two subterms that became different and were identical before; however, because the first step already took care of recipes introducing the same random seed twice in the same context, and the protocols in \mathcal{C}_{pp} cannot use a random seed from an input to use it in another encryption, it is impossible.

Hence, we showed that modifying tr into $\tilde{\text{tr}}$ is a symmetric operation that preserves equalities in the two protocols: identical plaintexts and keys in (tr, ϕ) correspond to identical plaintexts and keys in (tr, ψ) , whereas adding fresh nonces does not create any equality in $\tilde{\phi}$ or $\tilde{\psi}$. If (tr, ϕ) does not have any equivalent trace in Q , $(\tilde{\text{tr}}, \tilde{\phi})$ does not either. If there exists no frame ψ such that $(\text{tr}, \psi) \in \text{trace}(Q)$, then there will

exist no frame $\tilde{\psi}$ such that $(\tilde{\text{tr}}, \tilde{\psi}) \in \text{trace}(Q)$ as input filtering is not affected by our transformation. Otherwise, if there exists ψ such that $(\text{tr}, \psi) \in \text{trace}(Q)$ but ϕ and ψ are not statically equivalent, because our transformation preserves the terms in the frame, any pair of recipes that distinguishes between the two of them will distinguish $\tilde{\phi}$ and $\tilde{\psi}$. So we can always assume that the random seeds occurring in the recipes $f(u, v, w)$ in (tr, ϕ) are distinct.

Let us now define a corresponding trace $(\tilde{\text{tr}}, \tilde{\phi}) \in \text{trace}_{\text{fwd}}(\tilde{P})$:

- each sequence $\text{in}(c_i, R).\text{out}(c_i, w')$, where $c_i \neq c$, is left unchanged;
- each sequence $\text{in}(c, f(R, R_k, n)).\text{out}(c, w')$, where $R_k\phi \downarrow = k$ and $f \in \{\text{senc}, \text{aenc}, \text{sign}\}$, is replaced by $\text{in}(c_{k, \alpha(k)}^f, R).\text{out}(c_{k, \alpha(k)}^f, w')$;
- each sequence $\text{in}(c, g(R, R_k)).\text{out}(c, w')$, where $R_k\phi \downarrow = k$ and $g \in \{\text{sdec}, \text{adec}, \text{check}\}$, is replaced by $\text{in}(c_{k, \alpha(k)}^g, R).\text{out}(c_{k, \alpha(k)}^g, w')$.

Note that each recipe R and R_k is a variable w or the constant start . The corresponding frame $\tilde{\phi}$ is then defined according to our semantics. Since we have assumed that the random seeds occurring in the recipes in tr are distinct, we have that $\tilde{\phi} = \phi$.

Finally, because $(\text{tr}, \phi) \in \text{trace}^{\text{io}*}(P)$ has no equivalent in Q , and the definition of $(\tilde{\text{tr}}, \tilde{\phi})$ does not alter the filtering on inputs or equalities between terms in the frame, $(\tilde{\text{tr}}, \tilde{\phi}) \in \text{trace}_{\text{fwd}}(\tilde{P})$ has no equivalent in \tilde{Q} . \square

C. ENCODING A PROTOCOL INTO A REAL-TIME GPDA

C.1. Characterization of Trace Equivalence

LEMMA C.1. *Let P and Q be two protocols in \mathcal{C}_{pp} ; if $P \approx_{\text{fwd}} Q$, then for every trace $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ and every $w, w' \in \text{dom}(\sigma_P)$, if $w\sigma_P = w'\sigma_P = c$ for some constant c , then $w\sigma_Q = w'\sigma_Q = c'$ for some constant c' , where σ_Q is the frame such that $(\text{tr}, \sigma_Q) \in \text{trace}(Q)$.*

PROOF. First, note that the frame σ_Q mentioned in the lemma is unique up to some alpha-renaming of the randoms that occur in σ_P . Thus, the choice of the frame σ_Q does not change anything regarding the result that we want to prove.

Actually, the only nontrivial point to prove is that if $w\sigma_P = c$, then $w\sigma_Q$ is necessarily a constant too. Since protocols in \mathcal{C}_{pp} have a replication for every branch, consider the trace obtained by “playing twice” the trace tr in P and Q ; that is, given $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ with

$$\text{tr} = \text{in}(c_{i_1}, \text{start}).\text{out}(c_{i_1}, w_1) \dots \text{in}(c_{i_l}, w_{l-1}).\text{out}(c_{i_l}, w_l),$$

build $(\text{tr}', \sigma'_P) \in \text{trace}_{\text{fwd}}(P)$, where:

$$\begin{cases} \text{tr}' \stackrel{\text{def}}{=} \text{tr}.\tilde{\text{tr}} \\ \tilde{\text{tr}} \stackrel{\text{def}}{=} \text{in}(c_{i_1}, \text{start}).\text{out}(c_{i_1}, w_{|\phi|+1}) \dots \text{in}(c_{i_l}, w_{|\phi|+l}).\text{out}(c_{i_l}, w_{|\phi|+l}), \end{cases}$$

where every occurrence of start in tr is kept in $\tilde{\text{tr}}$ but occurrences of w_k are replaced by $w_{|\sigma_P|+k}$, $|\sigma_P|$ being the cardinal of $\text{dom}(\sigma_P)$; and $\text{tr}.\tilde{\text{tr}}$ denotes the concatenation of the two sequences of labels, which is a valid trace, that is, $(\text{tr}', \sigma'_P) \in \text{trace}_{\text{fwd}}(P)$. We get symmetrically $(\text{tr}', \sigma'_Q) \in \text{trace}_{\text{fwd}}(Q)$. In particular, there exists $w_* \in \text{dom}(\sigma'_P)$ with $l < *$ such that $w\sigma'_P = w_*\sigma'_P = c$ and the test $w = w_*$ is disjoint; that is, $\text{seq}_{\text{tr}}(w)$ and $\text{seq}_{\text{tr}'}(w_*)$ share no common prefix. As $P \approx_{\text{fwd}} Q$, necessarily $w\sigma'_Q = w_*\sigma'_Q$. Now, because the test is disjoint, $w\sigma'_Q$ and $w_*\sigma'_Q$ could not share any random nonces. Hence, $w\sigma_Q$ is a constant. \square

LEMMA C.2. *Let P and Q be two protocols in \mathcal{C}_{pp} such that $P \approx_{\text{fwd}} Q$. For every trace $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$, every $w, w' \in \text{dom}(\sigma_P)$ such that the test $w = w'$ is σ_P -valid, σ_P -guarded, and pulled up in (tr, σ_P) , we have that $w = w'$ is σ_Q -valid, σ_Q -guarded, and pulled up in (tr, σ_Q) , where σ_Q is the frame such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$.*

PROOF. First, note that the frame σ_Q mentioned in the lemma is unique up to some alpha-renaming of the randoms that occur in σ_P . Thus, the choice of the frame σ_Q does not change anything regarding the result that we want to prove.

The only nontrivial point to prove is that if the test $w = w'$ is σ_P -valid, σ_P -guarded, and pulled up in (tr, σ_P) , then it is also σ_Q -guarded and pulled up in (tr, σ_Q) . Note that it is necessarily σ_Q -valid since $P \approx_{\text{fwd}} Q$. Actually, we can still assume that the test $w = w'$ is σ_Q -guarded (it would otherwise contradict Lemma C.1).

Let $\text{pref} = \text{io}(c_{i_0}, \text{start}, w_{j_0}) \dots \text{io}(c_{i_p}, w_{j_{p-1}}, w_{j_p})$ be the maximal common prefix of $\text{seq}_{\text{tr}}(w)$ and $\text{seq}_{\text{tr}}(w')$. Now, it remains to show that $w = w'$ is pulled up in (tr, σ_Q) ; that is, $w\sigma_Q$ does not occur as a subterm in $w_{j_{-1}}\sigma_Q, w_{j_0}\sigma_Q, \dots, w_{j_{p-1}}\sigma_Q$, where $w_{j_{-1}}\sigma_Q = \text{start}$.

Assume that this is not the case; we will show that there exists a trace $(\text{tr}^*, \sigma_Q^*) \in \text{trace}_{\text{fwd}}(Q)$, $w_*, w'_* \in \text{dom}(\sigma_Q^*)$ such that $w_*\sigma_Q^* = w'_*\sigma_Q^*$, whereas $w_*\sigma_P^* \neq w'_*\sigma_P^*$, where σ_P^* is the frame such that $(\text{tr}^*, \sigma_P^*) \in \text{trace}_{\text{fwd}}(P)$. Note that such a frame necessarily exists since otherwise it trivially contradicts our hypothesis.

Let $p' \in \{0, \dots, p-1\}$ be the smallest index such that $w\sigma_Q$ occurs as a subterm in $w_{j_{p'}}\sigma_Q$. We have that

$$\text{pref} = s_1.\text{io}(c_{i_{p'}}, w_{j_{p'-1}}, w_{j_{p'}}).s_2 \quad \text{and} \quad \begin{cases} \text{seq}_{\text{tr}}(w) = \text{pref}.s_3 \\ \text{seq}_{\text{tr}}(w') = \text{pref}.s'_3 \end{cases}$$

for some sequence s_1, s_2, s_3 , and s'_3 .

From these sequences, we can define $(\text{tr}^*, \sigma_Q^*)$ with $\text{tr}^* = \text{tr}.\bar{\text{tr}}$. Intuitively, the trace $\bar{\text{tr}}$ is obtained relying on the sequence of channels as indicated in the sequence $s_2.s'_3$ using systematically the last generated recipe to feed the following input, and $w_{j_{p'}}$ to start. More precisely, assuming that

$$\text{seq}_{\text{tr}}(w') = s_1.\text{io}(c_{i_{p'}}, w_{j_{p'-1}}, w_{j_{p'}}).\text{io}(c_{k_1}, w_{j_{p'}}, w_{l_1}).\text{io}(c_{k_2}, w_{l_1}, w_{l_2}) \dots \text{io}(c_{k_\ell}, w_{l_{\ell-1}}, w_{l_\ell}),$$

we have that

$$\bar{\text{tr}} = \text{io}(c_{k_1}, w_{j_{p'}}, w_{|\sigma_P|+1}).\text{io}(c_{k_2}, w_{|\sigma_P|+1}, w_{|\sigma_P|+2}) \dots \text{io}(c_{k_\ell}, w_{|\sigma_P|+\ell-1}, w_{|\sigma_P|+\ell})$$

and σ_Q^* defined as expected relying on our semantics. Let $w_* = w$ and $w'_* = w_{|\sigma_P|+\ell}$. We can now show the following:

- (1) *The test $w_* = w'_*$ is σ_Q^* -valid and σ_Q^* -guarded.* Indeed, by definition of tr^* , $w_*\sigma_Q^*$ and $w'_*\sigma_Q^*$ are already equal up to a renaming of random seeds, as the channel components of $\text{seq}_{\text{tr}}(w')$ and $\text{seq}_{\text{tr}^*}(w'_*)$ match. As $w_*\sigma_Q^* = w\sigma_Q = w'\sigma_Q$, $w_*\sigma_Q^*$ and $w'_*\sigma_Q^*$ are equal up to a renaming of their random seeds. Lastly, we have that $w_*\sigma_Q^*$ and $w'_*\sigma_Q^*$ are both subterms of $w_{j_{p'}}\sigma_Q^*$, and hence $w_*\sigma_Q^* = w'_*\sigma_Q^*$.
- (2) *The test $w_* = w'_*$ is pulled up in $(\text{tr}^*, \sigma_Q^*)$.* This is by construction of tr^* .

Finally, as $P \approx_{\text{fwd}} Q$, there exists σ_P^* such that $(\text{tr}^*, \sigma_P^*) \in \text{trace}_{\text{fwd}}(P)$. But now $w_* = w'_*$ is σ_Q^* -valid, σ_Q^* -guarded, and pulled up in $(\text{tr}^*, \sigma_Q^*)$. Moreover, we are now in a situation where the top-level random seeds of $w_*\sigma_P^*$ and $w'_*\sigma_P^*$ are generated outside the common prefix of $\text{seq}_{\text{tr}^*}(w_*)$ and $\text{seq}_{\text{tr}^*}(w'_*)$, and thus this implies that $w_*\sigma_P^* \neq w'_*\sigma_P^*$, contradicting the equivalence $P \approx_{\text{fwd}} Q$. \square

LEMMA C.3. *Let P and Q be two protocols in \mathcal{C}_{pp} , and then $P \approx_{\text{fwd}} Q$ if and only if the following four conditions are satisfied:*

- CONST_P: *For all $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$, there exists a frame σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$ and for every $w, w' \in \text{dom}(\sigma_P)$ and for every constant $c \in \Sigma_0 \cup \{\text{start}\}$, $w\sigma_P = w'\sigma_Q = c$ if and only if there exists a constant $c' \in \Sigma_0 \cup \{\text{start}\}$ such that $w\sigma_Q = w'\sigma_Q = c'$.*
- CONST_Q: *Similarly swapping the roles of P and Q .*
- GUARDED_P: *For all $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$, there exists a frame σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$ and every test that is σ_P -valid, σ_P -guarded, and pulled up in (tr, σ_P) is also σ_Q -valid, σ_Q -guarded, and pulled up in (tr, σ_Q) .*
- GUARDED_Q: *Similarly swapping the roles of P and Q .*

PROOF. We prove the two directions separately.

(\Rightarrow) This implication is a direct consequence of Lemma C.1 and Lemma C.2.

(\Leftarrow) Suppose that $P \not\approx_{\text{fwd}} Q$. This means that there exists for instance $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ such that either there exists no frame σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$, in which case conditions CONST_P and GUARDED_P fail, or σ_Q is indeed defined and there exists a test $w = w'$ such that $w\sigma_P = w'\sigma_P$ but $w\sigma_Q \neq w'\sigma_Q$ (or the converse). Let us assume that $w\sigma_P = w'\sigma_P$ but $w\sigma_Q \neq w'\sigma_Q$.

If $w\sigma_P = w'\sigma_P = c$ for some constant c , then condition CONST_P is false.

Otherwise, we have that the test $w = w'$ is σ_P -valid and σ_P -guarded. From tr and $w = w'$, we will build a new trace $(\text{tr}^*, \sigma_P^*)$ and a new test $w_* = w'_*$ that is σ_P^* -valid, σ_P^* -guarded, and also pulled up in $(\text{tr}^*, \sigma_P^*)$. Actually, we proceed as in the proof of the previous lemma.

Let $\text{pref} = \text{io}(c_{i_0}, \text{start}, w_{j_0}) \dots \text{io}(c_{i_{p'}}, w_{j_{p'-1}}, w_{j_{p'}})$ be the maximal common prefix of $\text{seq}_{\text{tr}}(w)$ and $\text{seq}_{\text{tr}}(w')$. Let $p' \in \{0, \dots, p-1\}$ be the smallest index such that $w\sigma_P$ occurs as a subterm in $w_{j_{p'}}\sigma_P$. Note that if this index does not exist, then the test $w = w'$ is already pulled up in (tr, σ_P) and we are done.

We have that

$$\text{pref} = s_1.\text{io}(c_{i_{p'}}, w_{j_{p'-1}}, w_{j_{p'}}).s_2 \quad \text{and} \quad \begin{cases} \text{seq}_{\text{tr}}(w) = \text{pref}.s_3 \\ \text{seq}_{\text{tr}}(w') = \text{pref}.s'_3 \end{cases}$$

for some sequence s_1, s_2, s_3 , and s'_3 .

From these sequences, we can define $(\text{tr}^*, \sigma_P^*)$ with $\text{tr}^* = \text{tr}, \bar{\text{tr}}$. Intuitively, the trace $\bar{\text{tr}}$ is obtained relying on the sequence of channels as indicated in the sequence $s_2.s'_3$ using systematically the last-generated recipe to feed the following input, and $w_{j_{p'}}$ to start. More precisely, assuming that

$$\text{seq}_{\text{tr}}(w') = s_1.\text{io}(c_{i_{p'}}, w_{j_{p'-1}}, w_{j_{p'}}).\text{io}(c_{k_1}, w_{j_{p'}}, w_{l_1}).\text{io}(c_{k_2}, w_{l_1}, w_{l_2}) \dots \text{io}(c_{k_\ell}, w_{l_{\ell-1}}, w_{l_\ell}),$$

we have that

$$\bar{\text{tr}} = \text{io}(c_{k_1}, w_{j_{p'}}, w_{|\sigma_P|+1}).\text{io}(c_{k_2}, w_{|\sigma_P|+1}, w_{|\sigma_P|+2}) \dots \text{io}(c_{k_\ell}, w_{|\sigma_P|+\ell-1}, w_{|\sigma_P|+\ell})$$

and σ_P^* defined as expected relying on our semantics. Let $w_* = w$ and $w'_* = w_{|\sigma_P|+\ell}$.

Now, either there exists no frame σ_Q^* such that $(\text{tr}^*, \sigma_Q^*) \in \text{trace}(Q)$, in which case condition GUARDED_P fails obviously, or such a frame exists. In this case, by construction of tr^* , we have that the test $w_* = w'_*$ is σ_P^* -valid, σ_P^* -guarded, and pulled up in $(\text{tr}^*, \sigma_P^*)$.

In order to conclude, it remains to show that $w_*\sigma_Q^* \neq w'_*\sigma_Q^*$. We already know that $w\sigma_Q = w_*\sigma_Q^*$. Suppose *ad absurdum* that $w_*\sigma_Q^* = w'_*\sigma_Q^*$. Because the sequences of channels that occur in $\text{seq}_{\text{tr}}(w')$ and $\text{seq}_{\text{tr}^*}(w'_*)$ are the same, $w'\sigma_Q$ and $w'_*\sigma_Q^*$ are either constant and equal or of the form $f(u, k, r)$ with $f \in \{\text{senc}, \text{aenc}, \text{sign}\}$ and equal up

to a renaming of their random seeds. In the first case, it is enough to conclude that $w\sigma_Q = w'\sigma_Q$, which is absurd. In the second case, $w_*\sigma_Q^*$ and $w'_*\sigma_Q^*$, being randomized, must have equal top-level random seeds, implying that this nonce was introduced before $\text{io}(c_{i_{p'}}, w_{j_{p'-1}}, w_{j_{p'}})$ in the common prefix of their respective sequences. As the said prefix is also common to w and w' in tr , $w\sigma_Q$ and $w'\sigma_Q$ share the same top-level random seed and are thus equal, contradicting our hypothesis. Therefore: $w_*\sigma_Q^* \neq w'_*\sigma_Q^*$. Hence, GUARDED_P is false.

Finally, if $w\sigma_Q = w'\sigma_Q$ but $w\sigma_P \neq w'\sigma_P$, conditions CONST_Q and GUARDED_Q will similarly fail. \square

C.2. From Trace Equivalence to Language Equivalence

LEMMA C.4. *Let P and Q be two protocols in \mathcal{C}_{pp} ; the two real-time GPDA $\mathcal{A}_{\text{CONST}}^P$ and $\mathcal{A}_{\text{CONST}}^Q$ are such that*

$$P \text{ and } Q \text{ satisfy conditions } \text{CONST}_P \text{ and } \text{CONST}_Q \text{ if and only if } \mathcal{L}(\mathcal{A}_{\text{CONST}}^P) = \mathcal{L}(\mathcal{A}_{\text{CONST}}^Q).$$

PROOF. We prove the two implications separately.

(\Rightarrow) Assume that $\mathcal{L}(\mathcal{A}_{\text{CONST}}^P) \neq \mathcal{L}(\mathcal{A}_{\text{CONST}}^Q)$, and consider w.l.o.g. a word $u \in \mathcal{L}(\mathcal{A}_{\text{CONST}}^P) \setminus \mathcal{L}(\mathcal{A}_{\text{CONST}}^Q)$. We distinguish two cases depending on whether u is accepted in state q_0 or q_f .

Case $u = c_{i_1}.c_{i_2} \dots c_{i_l}$ is accepted in q_0 : In such a case, we built (tr, σ_P) as follows:

$$\text{tr} = \text{io}(c_{i_1}, \text{start}, w_1).\text{io}(c_{i_2}, w_1, w_2) \dots \text{io}(c_{i_l}, w_{l-1}, w_l),$$

with σ_P the substitution defined uniquely as expected from our semantics. We have that $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ as the transition function δ fully captures input filtering and output of terms for protocols in \mathcal{C}_{pp} . Since $u \notin \mathcal{L}(\mathcal{A}_{\text{CONST}}^Q)$, we have that $(\text{tr}, \sigma_Q) \notin \text{trace}_{\text{fwd}}(Q)$ for any substitution σ_Q , and thus the condition CONST_P fails.

Case u is accepted in q_f : In such a case, we also build a trace $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ “corresponding” to u . The construction is a bit more involved. We have that u is of the form $c_{i_1}.c_{i_2} \dots c_{i_k}.\text{ctest}c_{j_1}c_{j_2} \dots c_{j_l}.\text{cend}$. Let $\text{tr} = \text{tr}_1.\text{tr}_2$, with tr_1 and tr_2 defined as follows:

$$\begin{aligned} \text{---tr}_1 &= \text{io}(c_{i_1}, \text{start}, w_1).\text{io}(c_{i_2}, w_1, w_2) \dots \text{io}(c_{i_k}, w_{k-1}, w_k); \\ \text{---tr}_2 &= \text{io}(c_{j_1}, \text{start}, w_{k+1}).\text{io}(c_{j_2}, w_{k+1}, w_{k+2}) \dots \text{io}(c_{j_l}, w_{k+l-1}, w_{k+l}); \end{aligned}$$

and σ_P is defined uniquely as expected from our semantics, as P is deterministic. We have that $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ as the transition function δ fully captures input filtering and output of terms for protocols in \mathcal{C}_{pp} . We can now define $w = w_k$ and $w' = w_{k+l}$. Because the transitions from q_0 to q_c and then from q_c to q_f for some constant c were possible, we get that $w\sigma_P = w'\sigma_P = c$.

We know that $u = u_1.\text{ctest}u_2.\text{cend} \notin \mathcal{L}(\mathcal{A}_{\text{CONST}}^Q)$, and we may assume that u_1 and u_2 are both in $\mathcal{L}(\mathcal{A}_{\text{CONST}}^Q)$. Indeed, otherwise, this means that there exists no substitution σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$, and thus CONST_P fails, and the result holds. From now on, we assume that there exists σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$.

Now, let $q \xrightarrow{c;\alpha/\beta} q'$ be the first failing transition in the run of u in $\mathcal{A}_{\text{CONST}}^Q$. We distinguish several cases:

- (1) *Case $q = q_0$ and $q' = q_c$ for some constant c .* In such a case, $w\sigma_Q \neq c$ for any constant c , and $w\sigma_Q$ is thus a guarded term. The condition CONST_P fails.

- (2) *Case $q = q_c$ and $q' = q_t$ for some constant c .* In such a case, $w\sigma_Q = c$ but $w\sigma \neq c$, making CONST_P fail once again.

Hence, P and Q do not satisfy CONST_P . Symmetrically, if $u \in \mathcal{L}(\mathcal{A}_{\text{CONST}}^Q) \setminus \mathcal{L}(\mathcal{A}_{\text{CONST}}^P)$, the condition CONST_Q will fail.

(\Leftarrow) Assume that P and Q do not satisfy CONST_P (or CONST_Q), that is, there exists a trace $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ such that:

- (1) either there exists no σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$;
- (2) or there exist $w, w' \in \text{dom}(\sigma_P)$ and a constant c such that $w\sigma_P = w'\sigma_P = c$ but either $w\sigma_Q$ is not a constant or $w\sigma_Q$ is a constant but $w\sigma_Q \neq w'\sigma_Q$.

We consider such a trace of minimal length ℓ .

In the first case, thanks to minimality, we have that $\text{seq}_{\text{tr}}(w_\ell) = \text{tr}$. From tr , we build a word $u \in \mathcal{L}(\mathcal{A}_{\text{CONST}}^P)$ by extracting the channels that occur in tr , keeping the order. Since there does not exist σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$, and the transition function δ of the automaton fully captures input filtering and output of terms for protocols in \mathcal{C}_{pp} , we have that $u \notin \mathcal{L}(\mathcal{A}_{\text{CONST}}^Q)$.

In the second case, thanks to minimality, we have that tr is actually made up of all the actions that occur in $\text{seq}_{\text{tr}}(w)$ and $\text{seq}_{\text{tr}}(w')$ (note that these two sequences may share some actions). From tr , we built a word $u = u_1 c_{\text{test}} u_2 c_{\text{end}} \in \mathcal{L}(\mathcal{A}_{\text{CONST}}^P)$ as follows:

- u_1 is obtained by extracting the channels that occur in $\text{seq}_{\text{tr}}(w)$ preserving the order; and
- u_2 is obtained by extracting the channels that occur in $\text{seq}_{\text{tr}}(w')$ preserving the order.

As the transition function δ fully captures input filtering and output of terms for protocols in \mathcal{C}_{pp} , we get that upon reading c_{test} , $\mathcal{A}_{\text{CONST}}^P$ is in q_0 , and the transition $q_0 \xrightarrow{c_{\text{test}}; \omega_C / \omega} q_c$ is indeed possible as $w\sigma_P = c$; and similarly, upon reading the c_{end} , $\mathcal{A}_{\text{CONST}}^P$ is in q_c , and the transition $q_c \xrightarrow{c_{\text{end}}; \omega_C / \omega} q_f$ is indeed possible as $w'\sigma_P = c$, and hence $u \in \mathcal{L}(\mathcal{A}_{\text{CONST}}^P)$. What remains to show is that $u \notin \mathcal{L}(\mathcal{A}_{\text{CONST}}^Q)$. We distinguish two cases:

- Case $w\sigma_Q$ is not a constant.* In such a case, no transition $q_0 \xrightarrow{c_{\text{test}}; \omega_C / \omega} q_c$ will be possible after u_1 ; and
- Case $w\sigma_Q$ is a constant c but $w\sigma_Q \neq w'\sigma_Q$.* In such a case, the transition $q_0 \xrightarrow{c_{\text{end}}; \omega_C / \omega} q_c$ will not be possible after u_2 .

Hence, u cannot belong to $\mathcal{L}(\mathcal{A}_{\text{CONST}}^Q)$. This allows us to conclude. \square

LEMMA C.5. *Let P and Q be two protocols in \mathcal{C}_{pp} ; the two real-time GPDA $\mathcal{A}_{\text{GUARDED}}^P$ and $\mathcal{A}_{\text{GUARDED}}^Q$ are such that*

$$P \text{ and } Q \text{ satisfy conditions } \text{GUARDED}_P \text{ and } \text{GUARDED}_Q \text{ if and only if } \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^P) = \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q).$$

PROOF. We prove the two directions separately.

(\Rightarrow) Assume that $\mathcal{L}(\mathcal{A}_{\text{GUARDED}}^P) \neq \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q)$, and consider w.l.o.g. a word $u \in \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^P) \setminus \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q)$. We distinguish two cases depending on whether the word u is accepted in state q_0 or q_t .

Case $u = c_{i_1} c_{i_2} \dots c_{i_l}$ is accepted in q_0 : In such a case, we built (tr, σ_P) as follows:

$$\text{tr} = \text{io}(c_{i_1}, \text{start}, w_1). \text{io}(c_{i_2}, w_1, w_2) \dots \text{io}(c_{i_l}, w_{l-1}, w_l),$$

with σ_P the substitution defined uniquely as expected from our semantics. We have that $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ as the transition function δ fully captures input filtering and output of terms for protocols in \mathcal{C}_{pp} . Since $u \notin \mathcal{L}(\mathcal{A}_{\text{CONST}}^Q)$, we have that $(\text{tr}, \sigma_Q) \notin \text{trace}_{\text{fwd}}(Q)$ for any substitution σ_Q , and thus the condition GUARDED_P fails.

Case u is accepted in q_f : In such a case, we also build a trace $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ “corresponding” to u . The construction is a bit more involved. We have that u is of the form $c_{i_1}c_{i_2}\dots c_{i_k}c_{\text{fork}}^{i_0}c_{j_1}c_{j_2}\dots c_{j_l}c_{\text{test}}c_{p_1}c_{p_2}\dots c_{p_m}c_{\text{end}}$. Let $\text{tr} = \text{tr}_0.\text{tr}_1.\text{tr}_2$ with tr_0 , tr_1 , and tr_2 being defined as follows:

- $\text{tr}_0 = \text{io}(c_{i_1}, \text{start}, w_1).\text{io}(c_{i_2}, w_1, w_2)\dots\text{io}(c_{i_k}, w_{k-1}, w_k).\text{io}(c_{i_0}, w_k, w_{k+1})$;
- $\text{tr}_1 = \text{io}(c_{j_1}, w_{k+1}, w_{k+2}).\text{io}(c_{j_2}, w_{k+2}, w_{k+3})\dots\text{io}(c_{j_l}, w_{k+l}, w_{k+l+1})$; and
- $\text{tr}_2 = \text{io}(c_{p_1}, w_{k+l+1}, w_{k+l+2}).\text{io}(c_{p_2}, w_{k+l+2}, w_{k+l+3})\dots\text{io}(c_{p_m}, w_{k+l+m}, w_{k+l+m+1})$,

and σ_P is uniquely defined as expected from our semantics, as P is deterministic. We have that $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ as the transition function fully captures input filtering and output of terms for protocols in \mathcal{C}_{pp} . We can now define $w = w_{k+l+1}$ and $w' = w_{k+l+m+1}$. The test is σ_P -guarded (the index k associated to the stack symbol (fork, k) is indeed strictly positive), σ_P -valid (since the last transition from q_2 to q_f requires the stack to be identical to the stack before reading c_{test}), and pulled up in (tr, σ_P) (since the fork tiles allow us to control the first time the top-level random seed of w_{σ_P} appears in the frame).

We know that $u = u_0c_{\text{fork}}^{i_0}u_1c_{\text{test}}u_2c_{\text{end}} \notin \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q)$, and we may assume w.l.o.g. that $u_0c_{i_0}u_1$ and $u_0c_{i_0}u_2$ are both in $\mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q)$. Indeed, otherwise, this means that there exists no frame σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$, and thus GUARDED_P fails, and the result holds. From now on, we assume that there exists σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$.

Now, let $q \xrightarrow{c;\alpha/\beta} q'$ be the first failing transition in the run of u in $\mathcal{A}_{\text{GUARDED}}^Q$. We distinguish several cases:

- (1) *Case $q = q_0$ and $q' = q_1$.* Since we have already assumed that $u_0c_{i_0}u_1$ is in $\mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q)$, this means that the required transition does not exist because $\|v_i^j\| = 0$. In such a case, the test $w = w'$ (even if it was σ_Q -valid and σ_Q -guarded) cannot be a pulled-up one in (tr, σ_Q) . Thus, the condition GUARDED_P fails.
- (2) *Case $q = q_1$ and $q' = q_1$.* In such a case, this means that a fork tile cannot be unstacked, meaning that the corresponding test (even if it was σ_Q -valid and σ_Q -guarded) will not be pulled up in (tr, σ_Q) , and GUARDED_P is false.
- (3) *Case $q = q_1$ and $q' = q_2$.* In such a case, the problem occurs because the fork tile is not at the top of the stack upon becoming test. The corresponding test $w = w'$ will not be σ_Q -valid since w_{σ_Q} will contain a random seed that has been generated after the “forking point,” and thus this random seed cannot occur in w_{σ_Q} . Thus, the condition GUARDED_P fails.
- (4) *Case $q = q_2$ and $q' = q_f$.* In such a case, the test tile is not at the top of the stack upon reading the last letter of the word. The test is not σ_Q -valid. The stack at this point, without the test tile, is not identical to the stack before the fork tile turning test, making GUARDED_Q fail.

Hence, GUARDED_Q fails as soon as $u \notin \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q)$.

(\Leftarrow) Assume that P and Q do not satisfy conditions GUARDED_P (or GUARDED_Q), that is, there exists a trace $(\text{tr}, \sigma_P) \in \text{trace}_{\text{fwd}}(P)$ such that:

- (1) either there exists no σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$;
- (2) or (such a σ_Q exists) and there exists $w, w' \in \text{dom}(\sigma_P)$ such that the test $w = w'$ is σ_P -guarded and σ_P -valid, and pulled up in (tr, σ_P) , and

- either $w = w'$ is not σ_Q -valid,
- or $w = w'$ is not σ_Q -guarded,
- or $w = w'$ is not pulled-up in (tr, σ_Q) .

We consider such a trace of minimal length ℓ .

In the first case, thanks to the minimality, we have that $\text{seq}_{\text{tr}}(w_\ell) = \text{tr}$. From tr , we build a word $u \in \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^P)$ by extracting the channels that occur in tr , keeping the order. Since there does not exist σ_Q such that $(\text{tr}, \sigma_Q) \in \text{trace}_{\text{fwd}}(Q)$, and the transition function δ of the automaton fully captures input filtering and output of terms for protocols in \mathcal{C}_{pp} , we have that $u \notin \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q)$.

In the second case, thanks to minimality, we have that tr is actually made up of all the actions that occur in $\text{seq}_{\text{tr}}(w)$ and $\text{seq}_{\text{tr}}(w')$. These two sequences have a maximal common prefix pref that is not empty. Actually, we have that

$$\text{pref} = \text{io}(c_{i_1}, \text{start}, w_1). \text{io}(c_{i_2}, w_1, w_2) \dots \text{io}(c_{i_p}, w_{p-1}, w_p) \text{ for some } p \geq 1.$$

From tr , we build a word $u = c_{i_1}c_{i_2} \dots c_{i_{p-1}}c_{\text{fork}}^{i_p}u_1c_{\text{test}}u_2c_{\text{end}} \in \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^P)$ as follows:

- u_1 is obtained by extracting the channels that occur in $\text{seq}_{\text{tr}}(w)$ after the prefix $c_{i_1}c_{i_2} \dots c_{i_{p-1}}c_{i_p}$; and
- u_2 is obtained by extracting the channels that occur in $\text{seq}_{\text{tr}}(w')$ after the prefix $c_{i_1}c_{i_2} \dots c_{i_{p-1}}c_{i_p}$.

As the transition function δ fully captures input filtering and output of terms for protocols in \mathcal{C}_{pp} , and since $w = w'$ is a test that is σ_P -guarded, σ_P -valid, and pulled up in (tr, σ_P) , we get that $u \in \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^P)$. What remains to show is that $u \notin \mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q)$. We distinguish two cases:

- Case $w = w'$ is not σ_Q -valid.* In such a case, even if after reading the first part of u , that is, $c_{i_1}c_{i_2} \dots c_{i_{p-1}}c_{\text{fork}}^{i_p}u_1c_{\text{test}}$, we reach q_2 , then we will fail to read the remaining of the word to end in q_f .
- Case $w = w'$ is σ_Q -valid but $w = w'$ is not σ_Q -guarded.* In such a case, this means that w_{σ_Q} is a constant, and the run will stop in q_0 after reading $c_{i_1}c_{i_2} \dots c_{i_{p-1}}$. This comes from the fact that it is not possible to go from q_0 to q_1 adding a tile (fork_i^j, k) with $k = 0$.
- Case $w = w'$ is σ_Q -valid and σ_Q -guarded but not pulled up in (tr, σ_Q) .* The fact that the test is σ_Q -valid but not pulled up means that the run will stop in q_1 after reading because of the presence of a tile (fork_i^j, k) in the stack that cannot go down anymore.

Hence, u cannot belong to $\mathcal{L}(\mathcal{A}_{\text{GUARDED}}^Q)$. This allows us to conclude. \square

REFERENCES

- M. Arapinis, T. Chothia, E. Ritter, and M. Ryan. 2010. Analysing unlinkability and anonymity using the applied pi calculus. In *23rd Computer Security Foundations Symposium (CSF'10)*. IEEE Computer Society Press, 107–121.
- D. Basin, S. Mödersheim, and L. Viganò. 2005. A symbolic model checker for security protocols. *Journal of Information Security* 4, 3 (2005), 181–208.
- M. Baudet. 2005. Deciding security of protocols against off-line guessing attacks. In *12th ACM Conference on Computer and Communications Security (CCS'05)*. ACM Press. DOI: <http://dx.doi.org/10.1145/1102125>
- B. Blanchet. 2001. An efficient cryptographic protocol verifier based on prolog rules. In *14th Computer Security Foundations Workshop (CSFW'01)*. IEEE Computer Society Press.
- B. Blanchet, M. Abadi, and C. Fournet. 2005. Automated verification of selected equivalences for security protocols. In *20th Symposium on Logic in Computer Science*.

- S. Boehm and S. Goeller. 2011. Language equivalence of deterministic real-time one-counter automata is NL-complete. In *Proceedings of the 36th International Symposium on Mathematical Foundations of Computer Science (MFCS'11) (LNCS)*, Vol. 6907. 194–205.
- M. Bruso, K. Chatzikokolakis, and J. den Hartog. 2010. Formal verification of privacy for RFID systems. In *23rd Computer Security Foundations Symposium (CSF'10)*.
- V. Cheval and B. Blanchet. 2013. Proving more observational equivalences with ProVerif. In *2nd Conference on Principles of Security and Trust (POST'13) (Lecture Notes in Computer Science)*, David Basin and John Mitchell (Eds.), Vol. 7796. Springer, Rome, Italy, 226–246.
- V. Cheval, H. Comon-Lundh, and S. Delaune. 2011. Trace equivalence decision: Negative tests and non-determinism. In *18th ACM Conference on Computer and Communications Security (CCS'11)*. ACM Press. <http://www.lsv.ens-cachan.fr/Publis/PAPERS/PDF/CCD-ccs11.pdf>.
- Y. Chevalier and M. Rusinowitch. 2012. Decidability of equivalence of symbolic derivations. *Journal of Automated Reasoning* 48, 2 (2012), 263–292.
- H. Comon-Lundh and V. Cortier. 2003. New decidability results for fragments of first-order logic and application to cryptographic protocols. In *14th International Conference on Rewriting Techniques and Applications (RTA'03) (LNCS)*, Vol. 2706. Springer.
- V. Cortier and S. Delaune. 2009. A method for proving observational equivalence. In *22nd IEEE Computer Security Foundations Symposium (CSF'09)*. IEEE Computer Society Press. DOI:<http://dx.doi.org/10.1109/CSF.2009.9>
- C. Cremers. 2008. Unbounded verification, falsification, and characterization of security protocols by pattern refinement. In *15th ACM Conference on Computer and Communications Security (CCS'08)*. ACM. DOI:<http://dx.doi.org/10.1145/1455770.1455787>
- D. Denning and G. Sacco. 1981. Timestamps in key distribution protocols. *Communication of the ACM* 24, 8 (1981), 533–536.
- E. P. Friedman. 1976. The inclusion problem for simple languages. *Theoretical Computer Science* 1, 4 (1976), 297–316.
- P. Henry and G. Sénizergues. 2013. LALBLC a program testing the equivalence of dpda's. In *18th International Conference on Implementation and Application of Automata (CIAA'13) (Lecture Notes in Computer Science)*, Vol. 7982. Springer, Halifax, NS, Canada, 169–180.
- ICAO. 2008. *Machine Readable Travel Document*. Technical Report 9303. International Civil Aviation Organization.
- M. Rusinowitch and M. Turuani. 2003. Protocol insecurity with finite number of sessions and composed keys is NP-complete. *Theoretical Computer Science* 299 (April 2003), 451–475. <http://www.loria.fr/~rusi/pub/tcpsprotocol.ps.gz>.
- G. Sénizergues. 1997. The equivalence problem for deterministic pushdown automata is decidable. In *24th International Colloquium on Automata, Languages and Programming (ICALP'97) (Lecture Notes in Computer Science)*.
- G. Sénizergues. 2001. $L(A)=L(B)$? Decidability results from complete formal systems. *Theoretical Computer Science* 251, 1–2 (2001), 1–166.
- C. Stirling. 2002. Deciding DPDA equivalence is primitive recursive. In *29th International Colloquium on Automata, Languages and Programming (ICALP'02) (LNCS)*. Springer.
- A. Tiu and J. E. Dawson. 2010. Automating open bisimulation checking for the spi calculus. In *23rd IEEE Computer Security Foundations Symposium (CSF'10)*. 307–321.

Received June 2014; revised June 2015; accepted July 2015